

# Weed Classification

Group project for CS3244

## **Presented by PG21:**

Yun Jihyun

Lee Chaeyeon

N Mufiz Ahmed

Son Seongyeon

Low Jia Jun

# *Contents*

*01.*

Introduction

*02.*

Pre-processing

*03.*

Linear Models

*04.*

Linear Models  
with Autoencoders

*05.*

CNN and GradCAM

*06.*

Conclusion

01

# Introduction

# *Context*

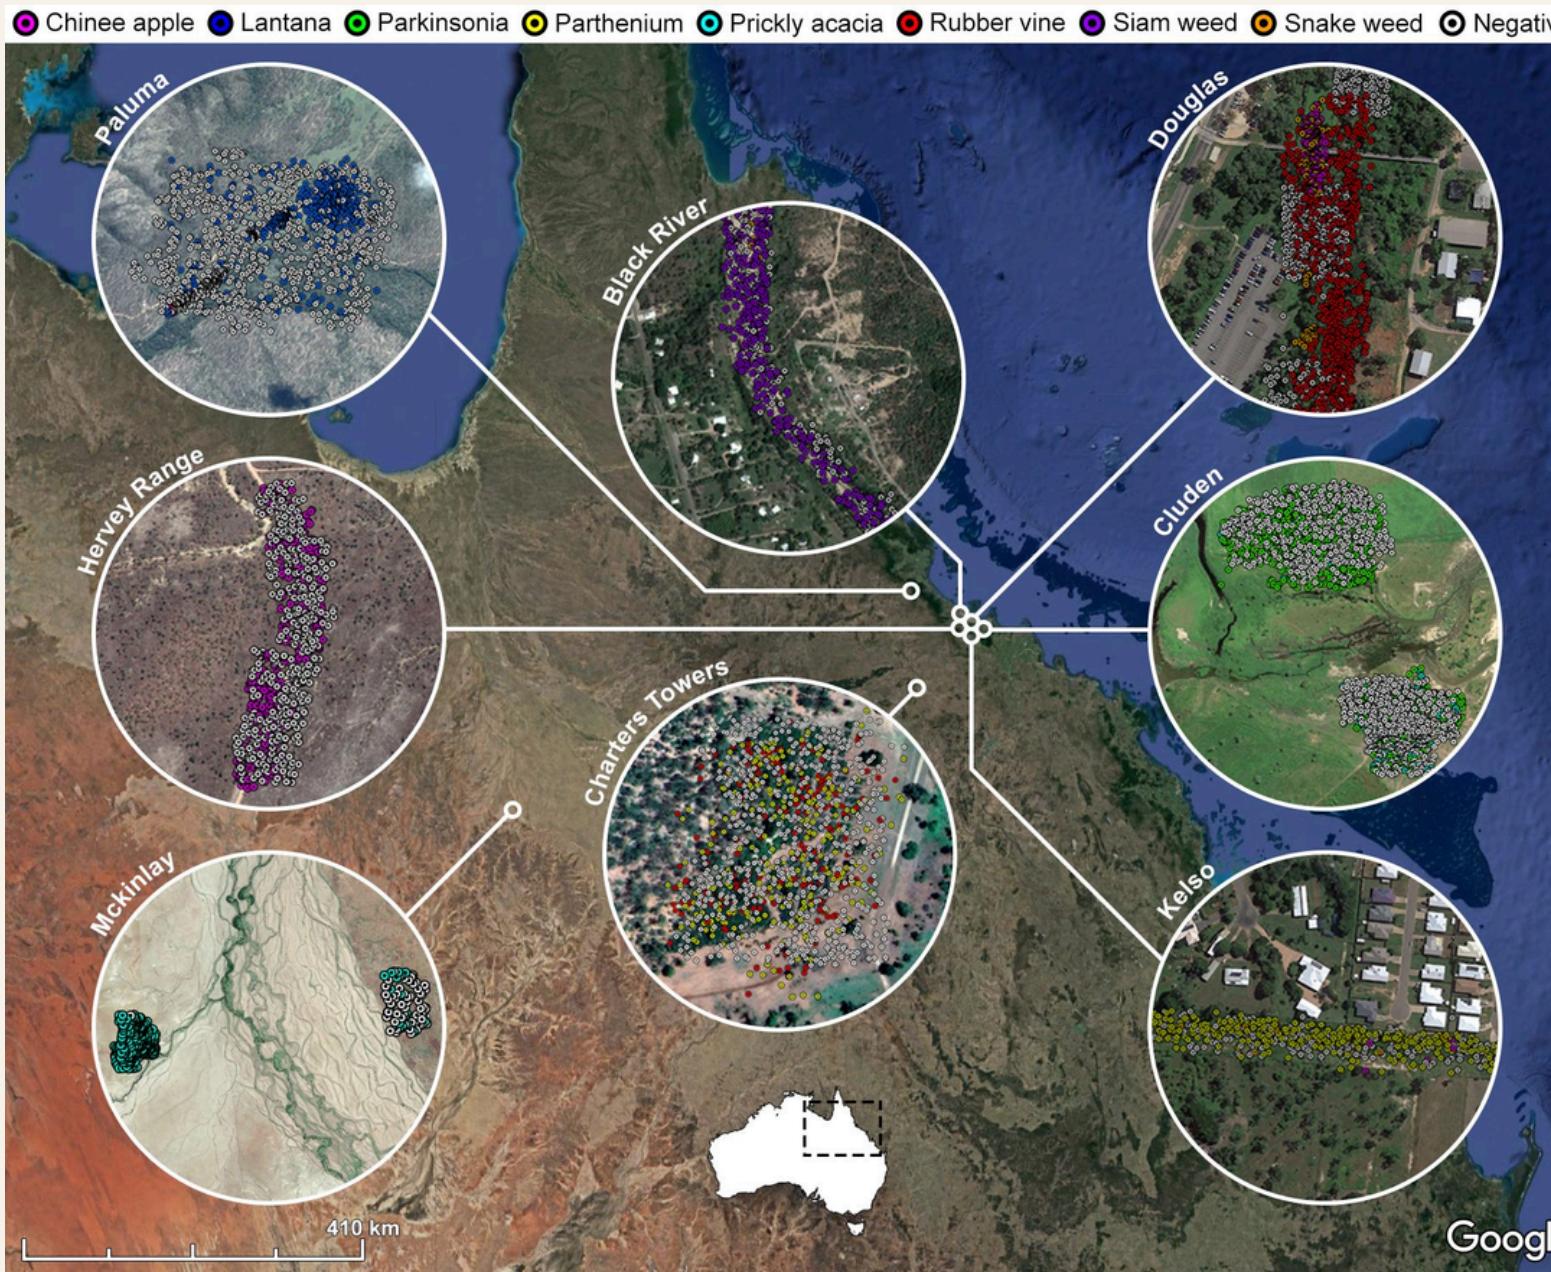
## **Economic**

Costing Australian farmers billions annually

## **Technological**

Growing need for intelligent systems

# Dataset Description



The dataset contains  
over 17500 images of  
8 different weed species  
native to Australia

# *Models used*

*01.*

Linear Models  
(kNN, Random  
Forest, SVM)

*02.*

Linear Models  
with  
Autoencoders

*03.*

Convolutional  
Neural Network

# *Motivation*

01.

Which model delivers the best accuracy for real-world weed classification?

02.

How do preprocessing techniques impact model performance?

# *Baseline Model*

Macro-averaged Precision: 0.06

Macro-averaged Recall: 0.11

Macro-averaged F1 score: 0.08

Classification Report for Logistic Regression:

	precision	recall	f1-score	support
Chinese Apple	0.00	0.00	0.00	213
Lantana	0.00	0.00	0.00	215
Parkinsonia	0.00	0.00	0.00	199
Parthenium	0.00	0.00	0.00	215
Prickly acacia	0.00	0.00	0.00	210
Rubber vine	0.00	0.00	0.00	200
Siam weed	0.00	0.00	0.00	223
Snake Weed	0.00	0.00	0.00	191
Negatives	0.52	1.00	0.69	1836
accuracy			0.52	3502
macro avg	0.06	0.11	0.08	3502
weighted avg	0.27	0.52	0.36	3502

Logistic Regression with macro-averaged F1 score of 0.08

# 02

# Data Preprocessing

# *Preprocessing*

01.

Inspection and class imbalance

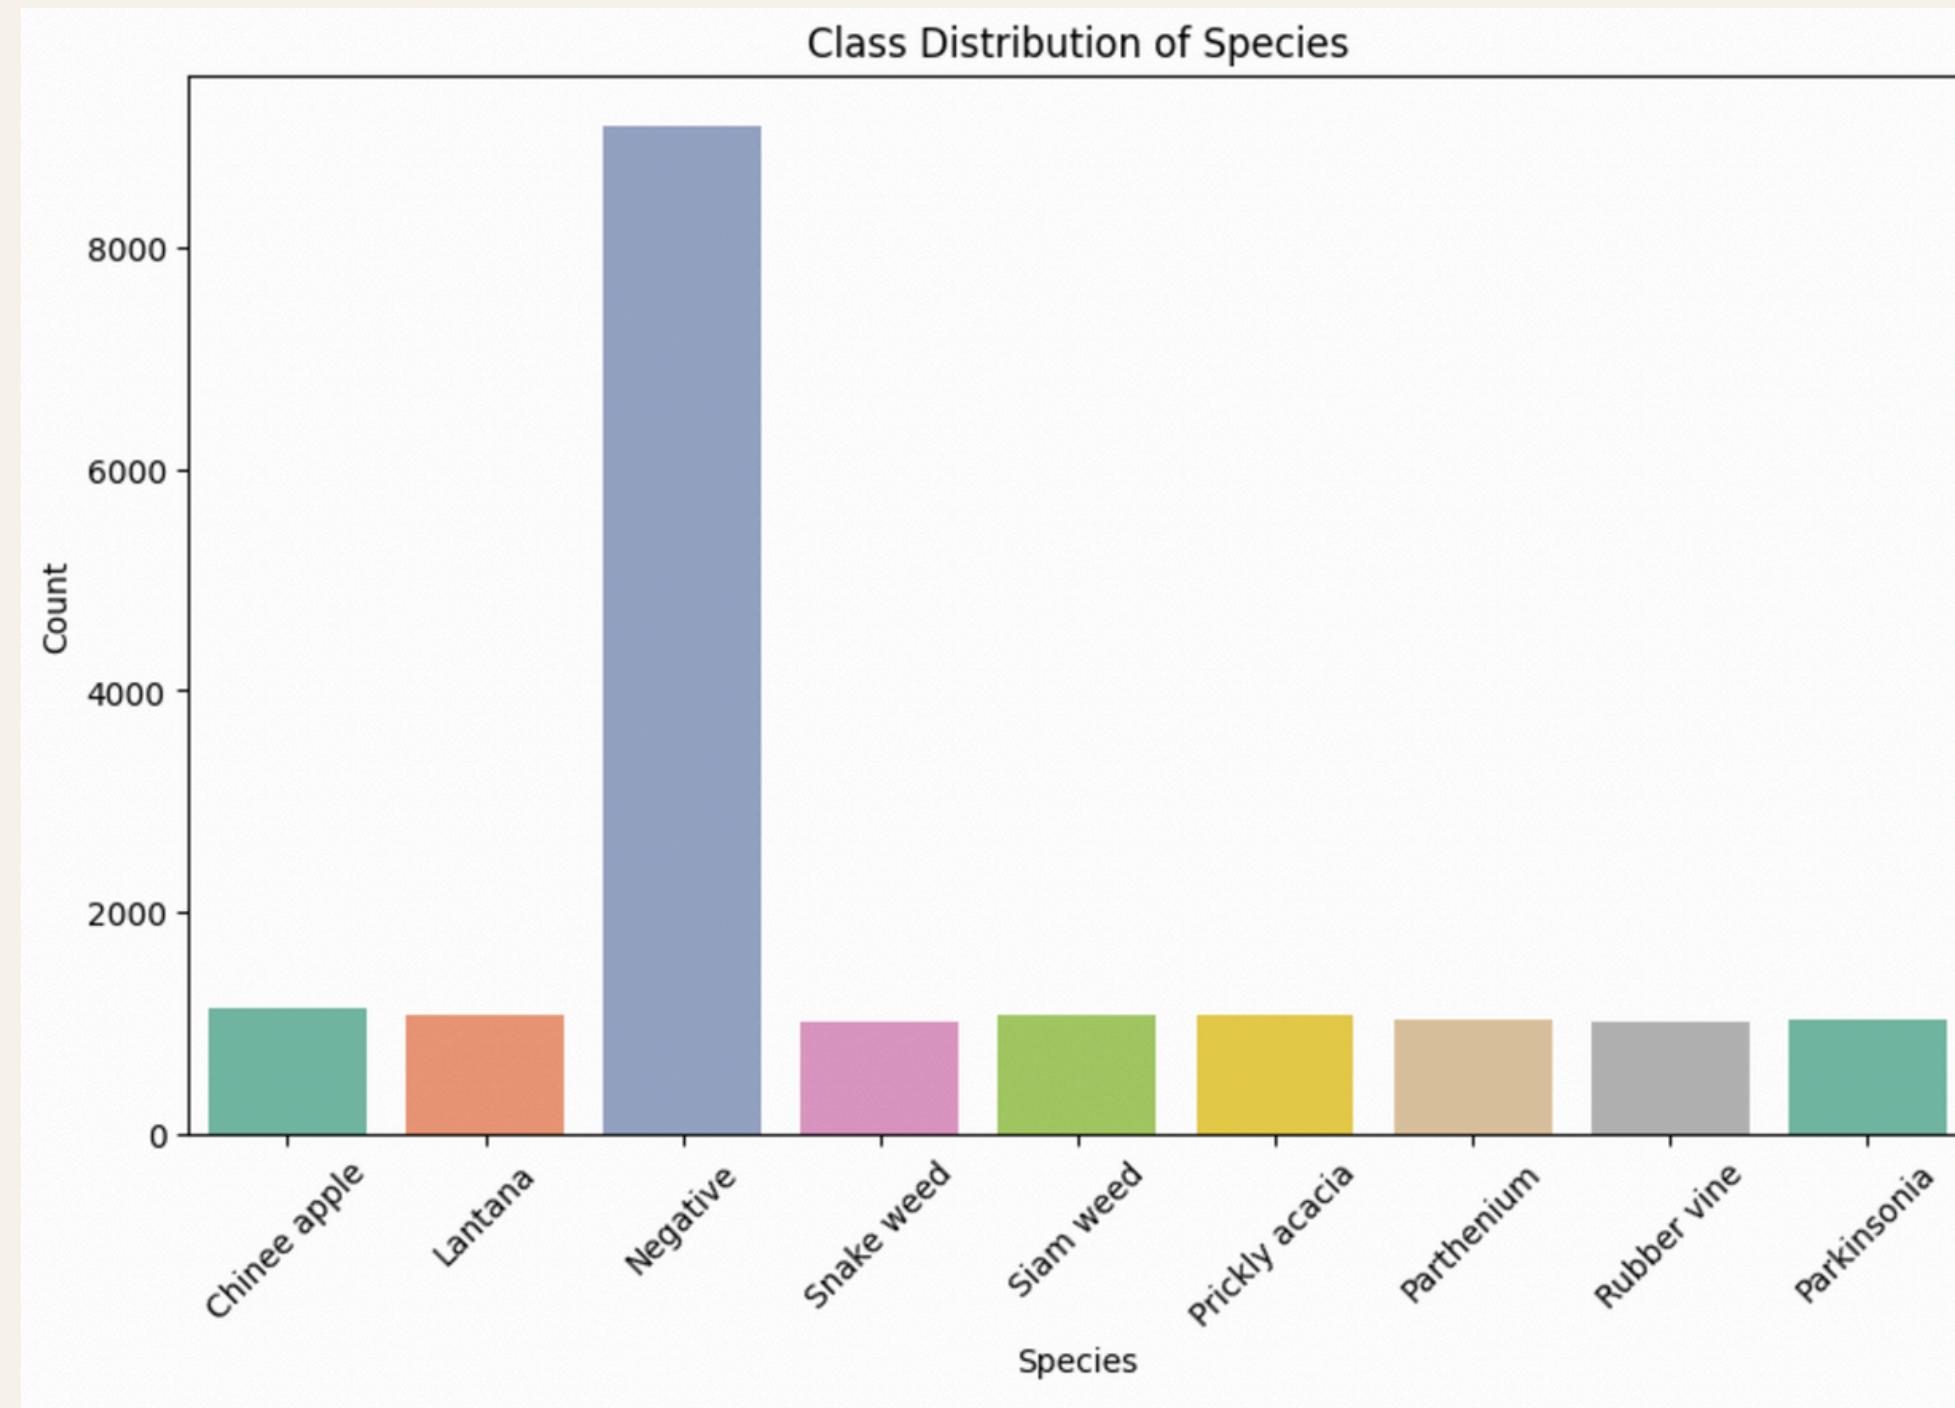
02.

Splitting and preprocessing

03.

Feature extraction

# *Inspection and class imbalance*



Negative class  
is overwhelmingly  
dominant which will  
reduce its ability  
to accurately classify  
other weed species

# *Inspection and class imbalance*

## **Oversampling?**

will lead to increased  
dataset size and potentially  
cause overfitting



## **Undersampling**

Randomly selected a subset  
of the negative class.  
Opted for this and achieved  
a balanced dataset



# *Splitting and Feature Extraction*

## Data splitting

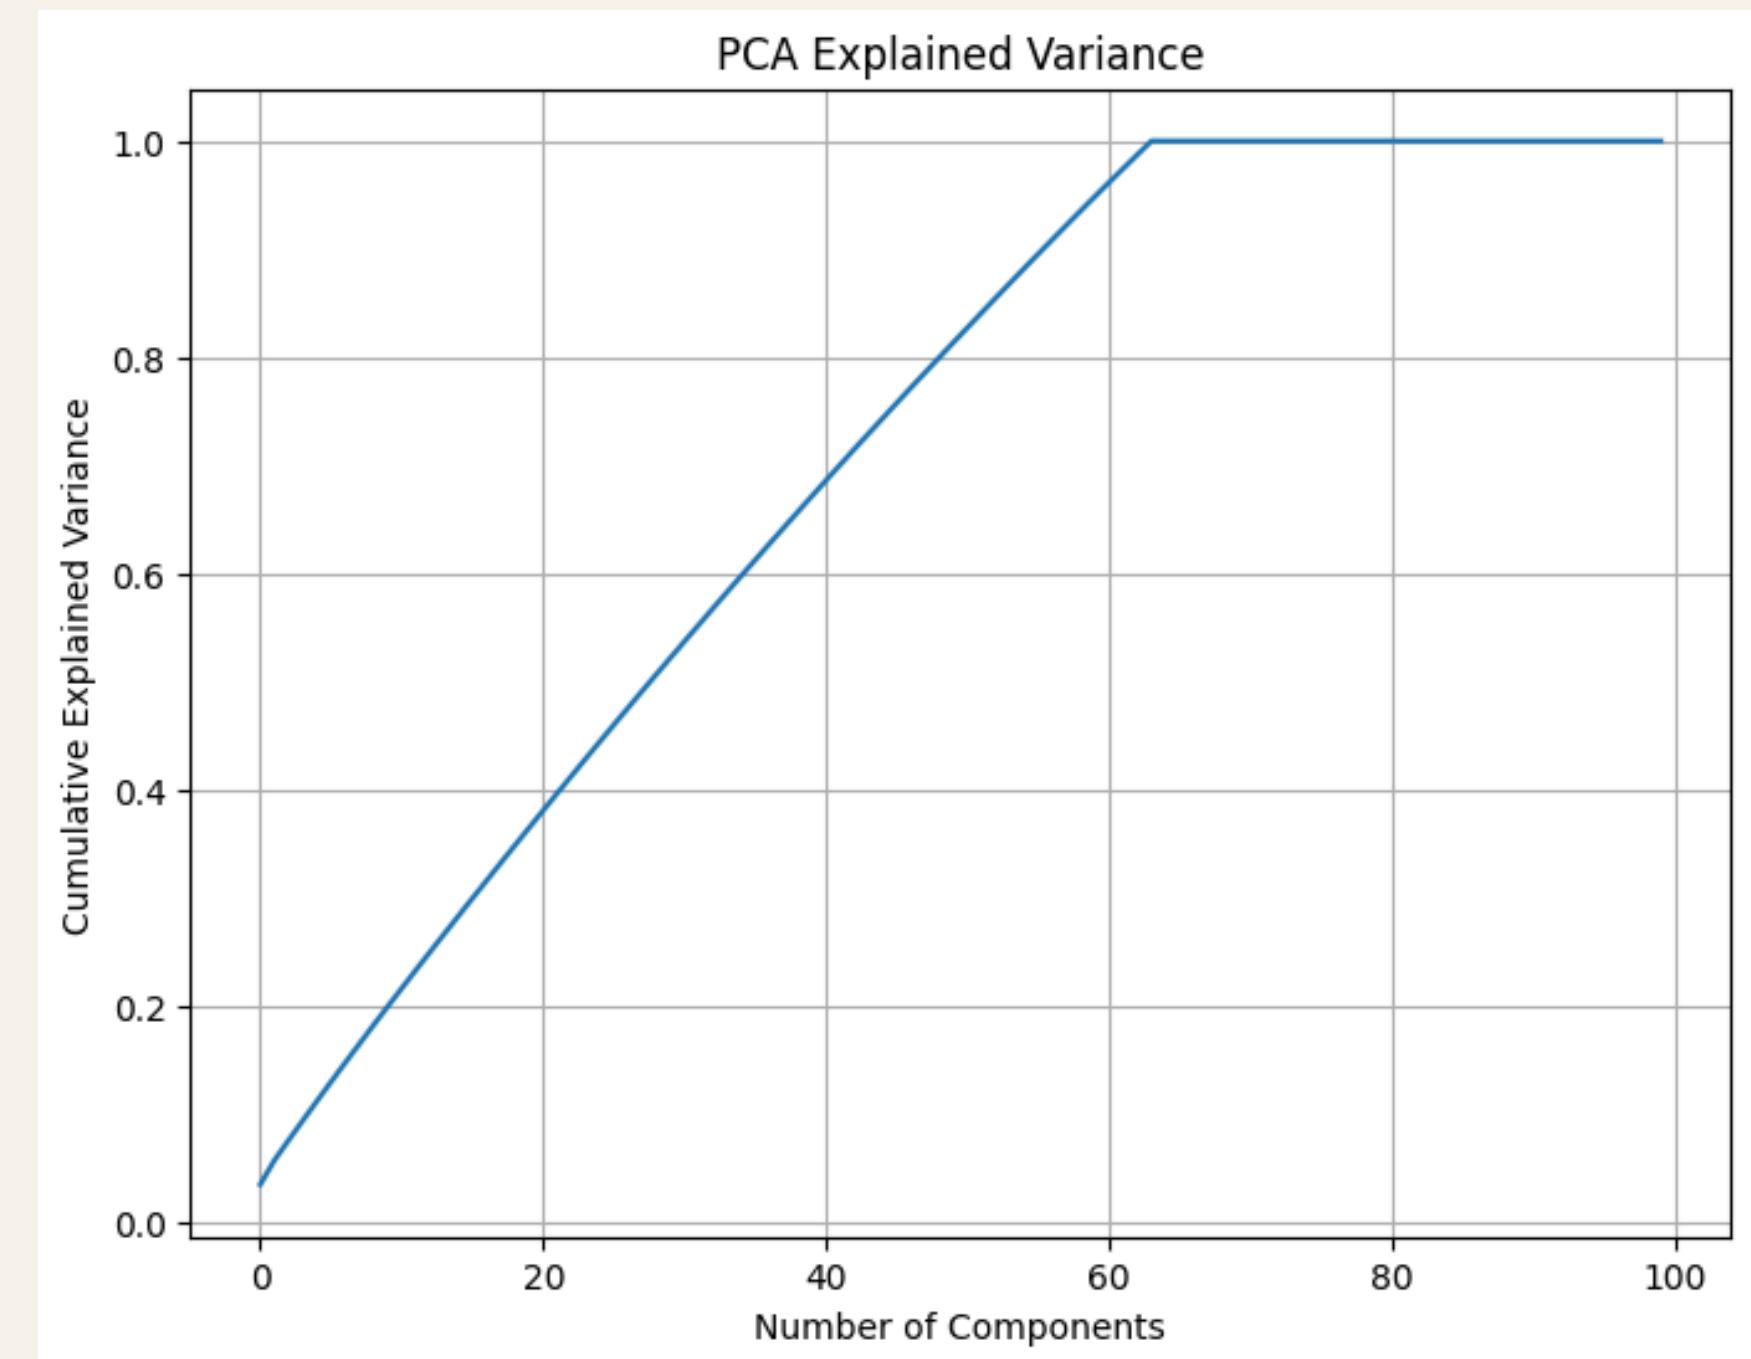
- Split the data into training and test sets to prevent data leakage
- Resized images to 224 x 224 pixels and normalized pixel values to [0,1]

## Features extracted

- Mean RGB Values: Capture overall color, key for distinguishing weed species.
- Local Binary Patterns (LBP): Encode texture, useful for species classification.
- Laplacian Mean Value: Highlights sharp intensity changes, emphasizing leaf edges.
- Edge Mean Value: Quantifies edge intensity, reinforcing structural features.
- Harris Corner Mean Value: Detects key points and corners, capturing unique weed shapes.

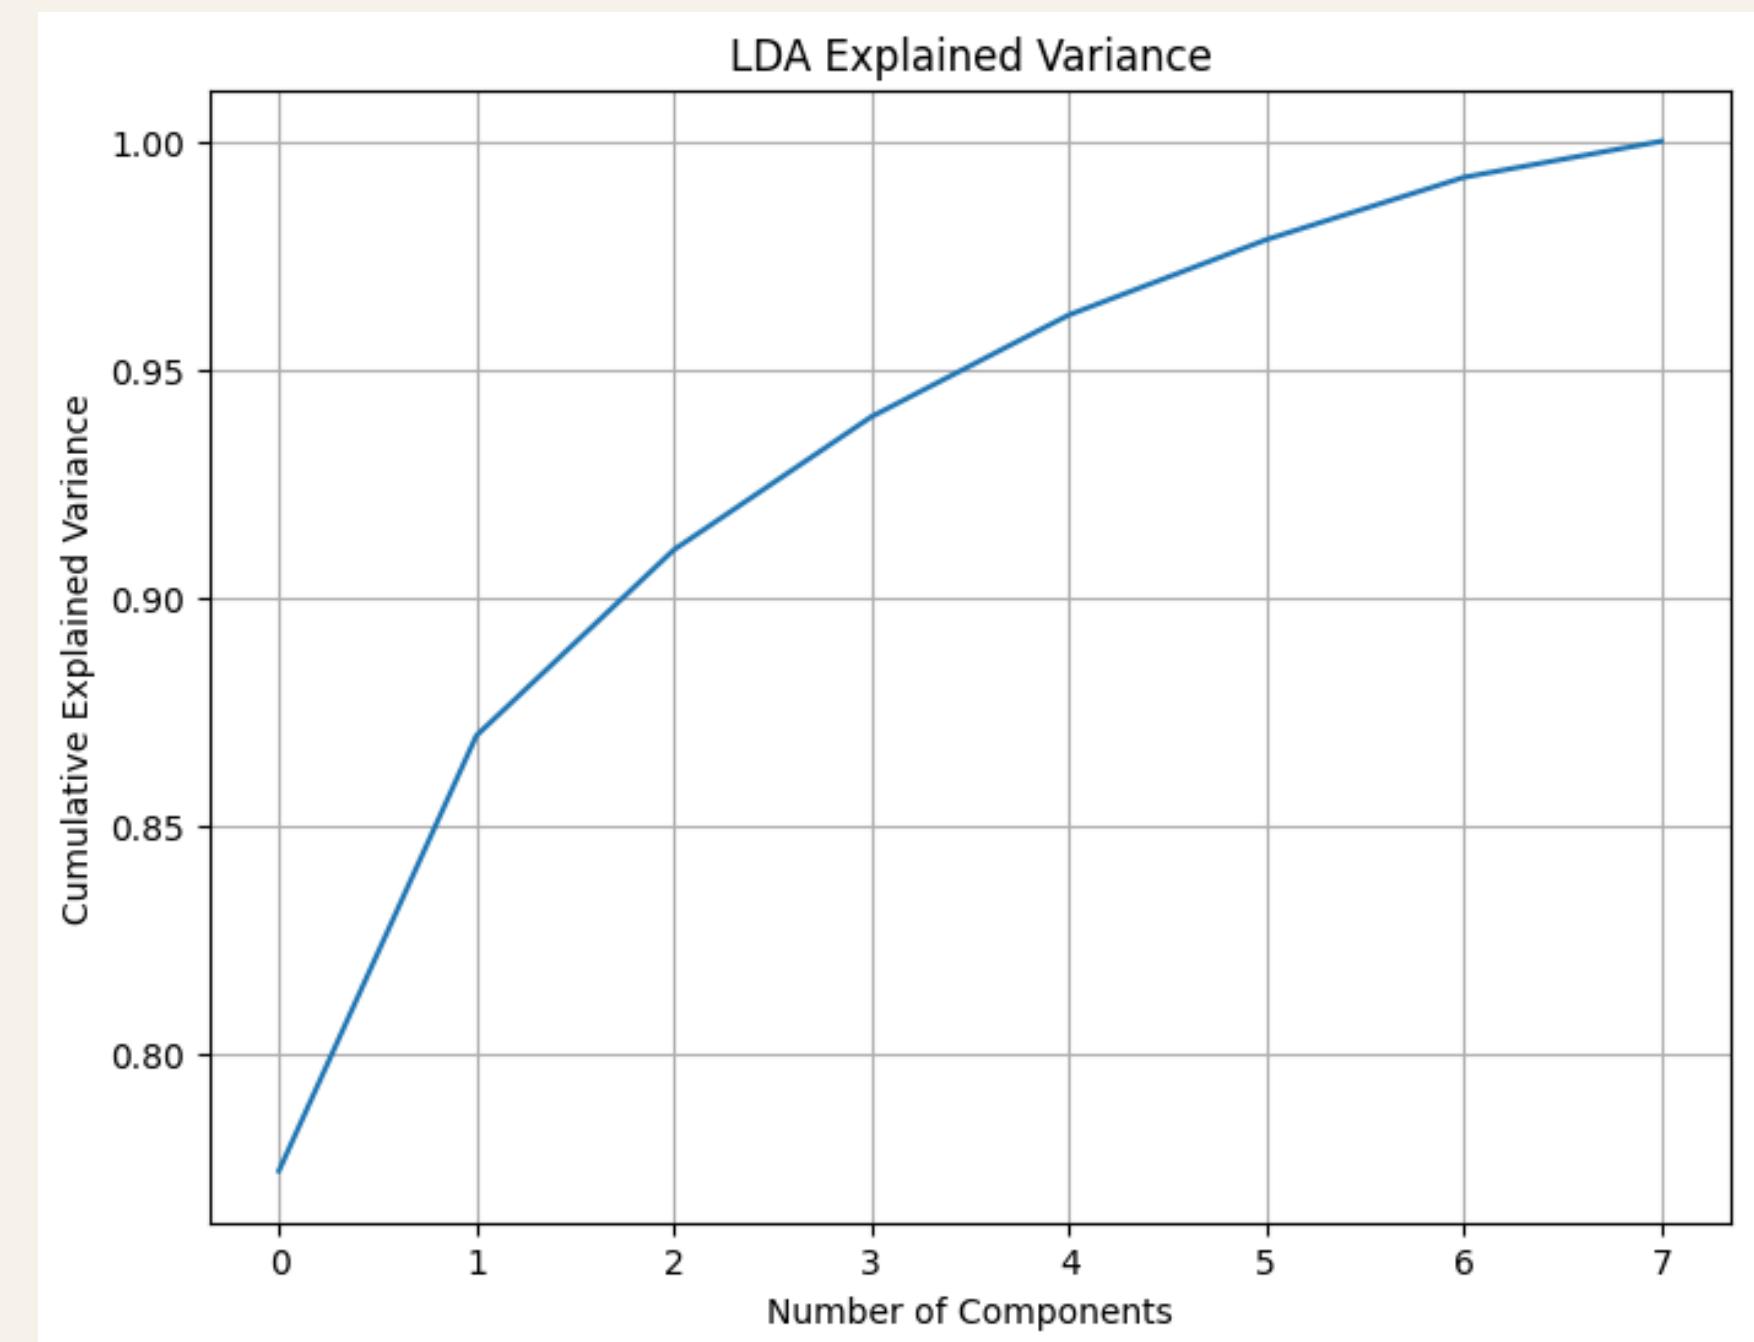
# *Dimensionality reduction (PCA)*

**150,535** features  
reduced to **64**!



# *Dimensionality reduction (LDA)*

**64** features  
reduced to **8!!**



# 03

# Linear Models

# *Linear Models*

01.

k-Nearest Neighbours (kNN)

02.

Random Forest Classifier

03.

Support Vector Machine (SVM)

# *k*-Nearest Neighbours (*kNN*)

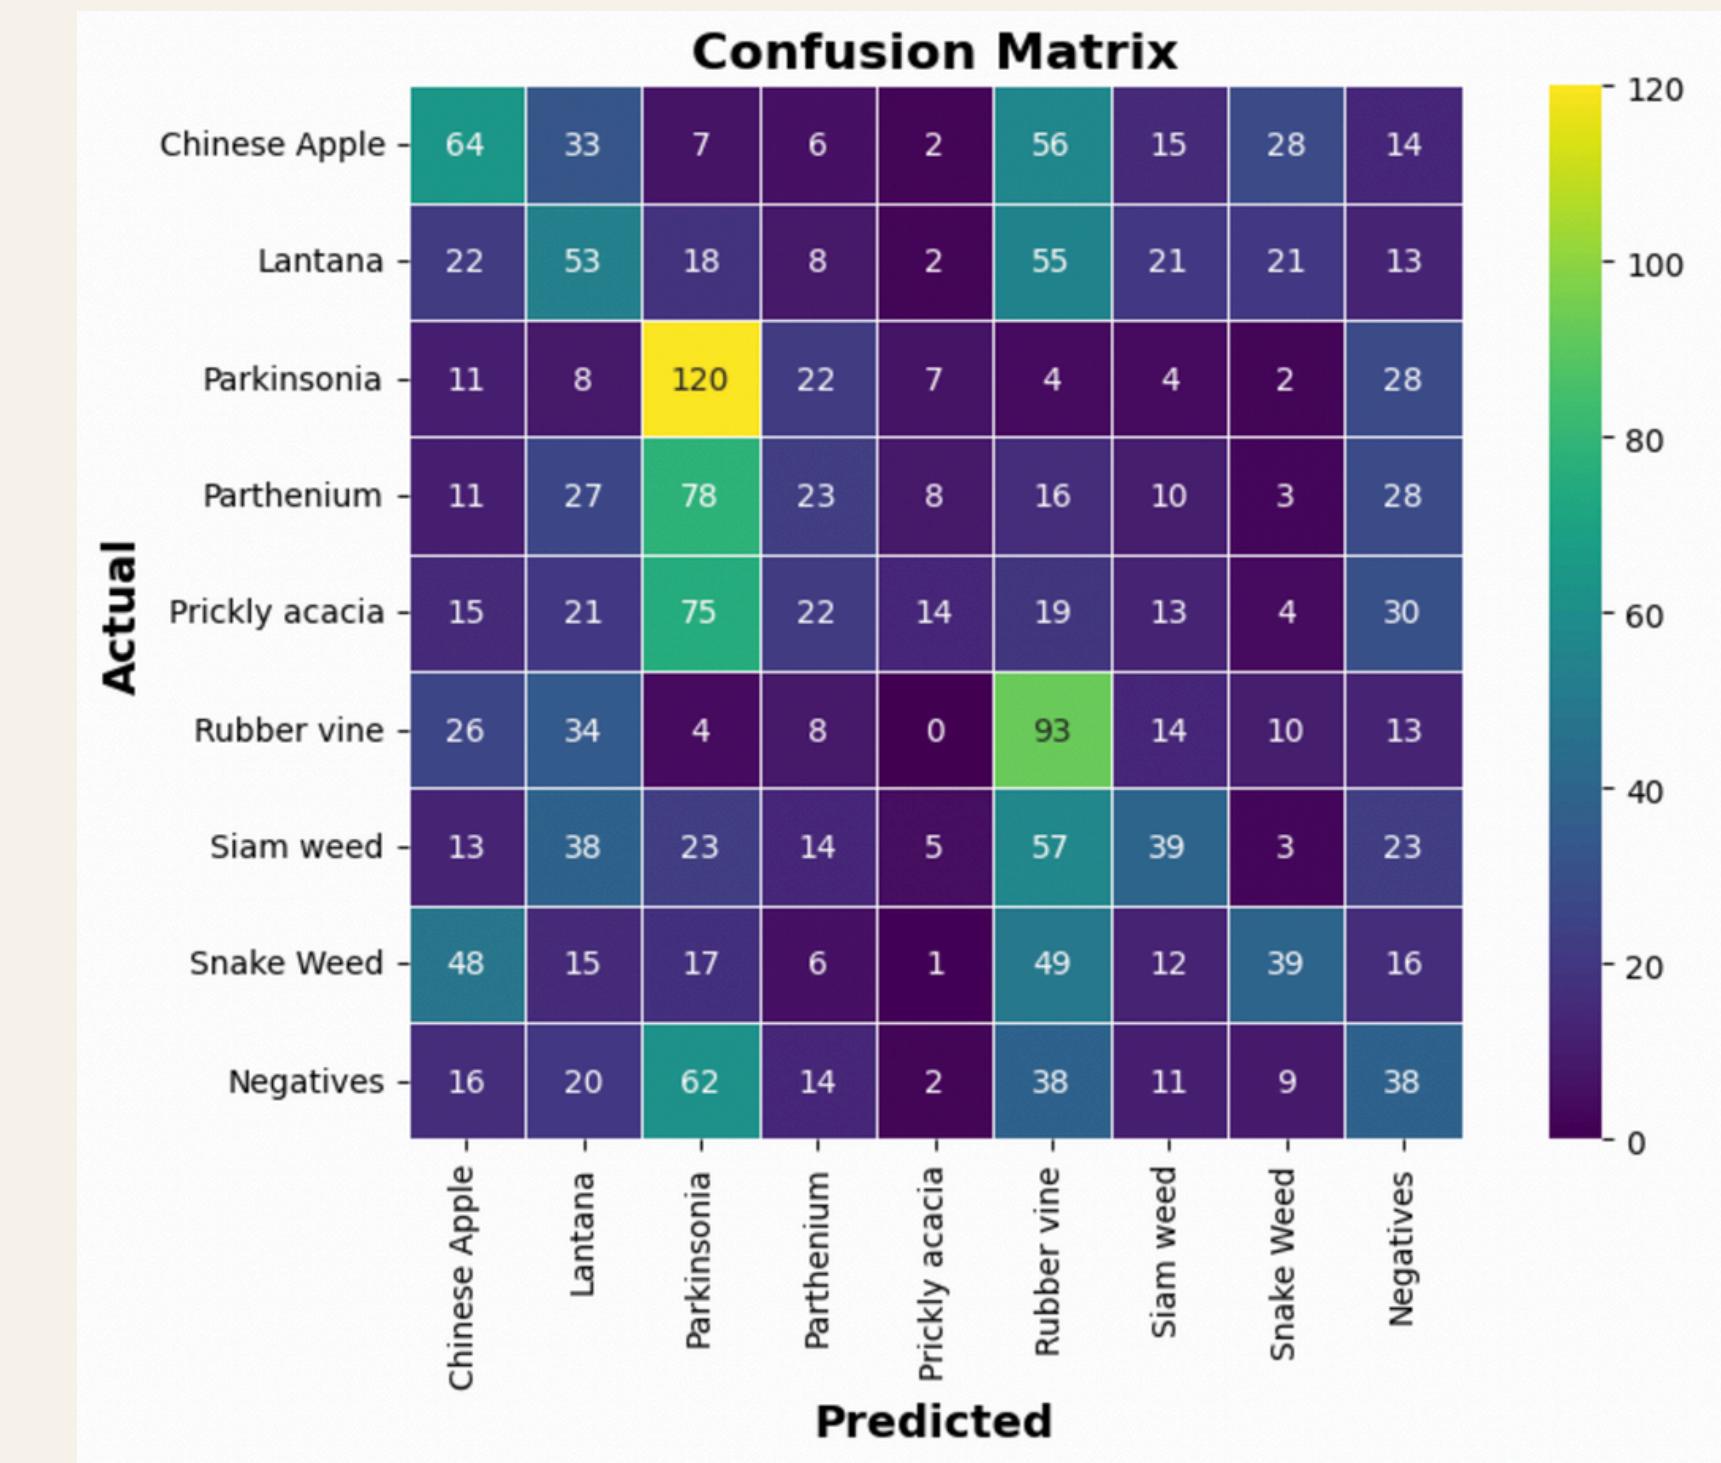
Micro-averaged Precision: 0.22

Micro-averaged Recall: 0.22

Micro-averaged F1 score: 0.22

Classification Report for kNN:

	precision	recall	f1-score	support
Chinese Apple	0.31	0.22	0.26	225
Lantana	0.15	0.13	0.14	213
Parkinsonia	0.13	0.75	0.22	206
Parthenium	0.14	0.14	0.14	204
Prickly acacia	0.00	0.00	0.00	213
Rubber vine	0.13	0.57	0.21	202
Siam weed	0.25	0.16	0.20	215
Snake Weed	0.14	0.13	0.13	203
Negatives	0.66	0.18	0.29	1821
accuracy			0.22	3502
macro avg	0.21	0.25	0.18	3502
weighted avg	0.42	0.22	0.23	3502



# Random Forest Classifier

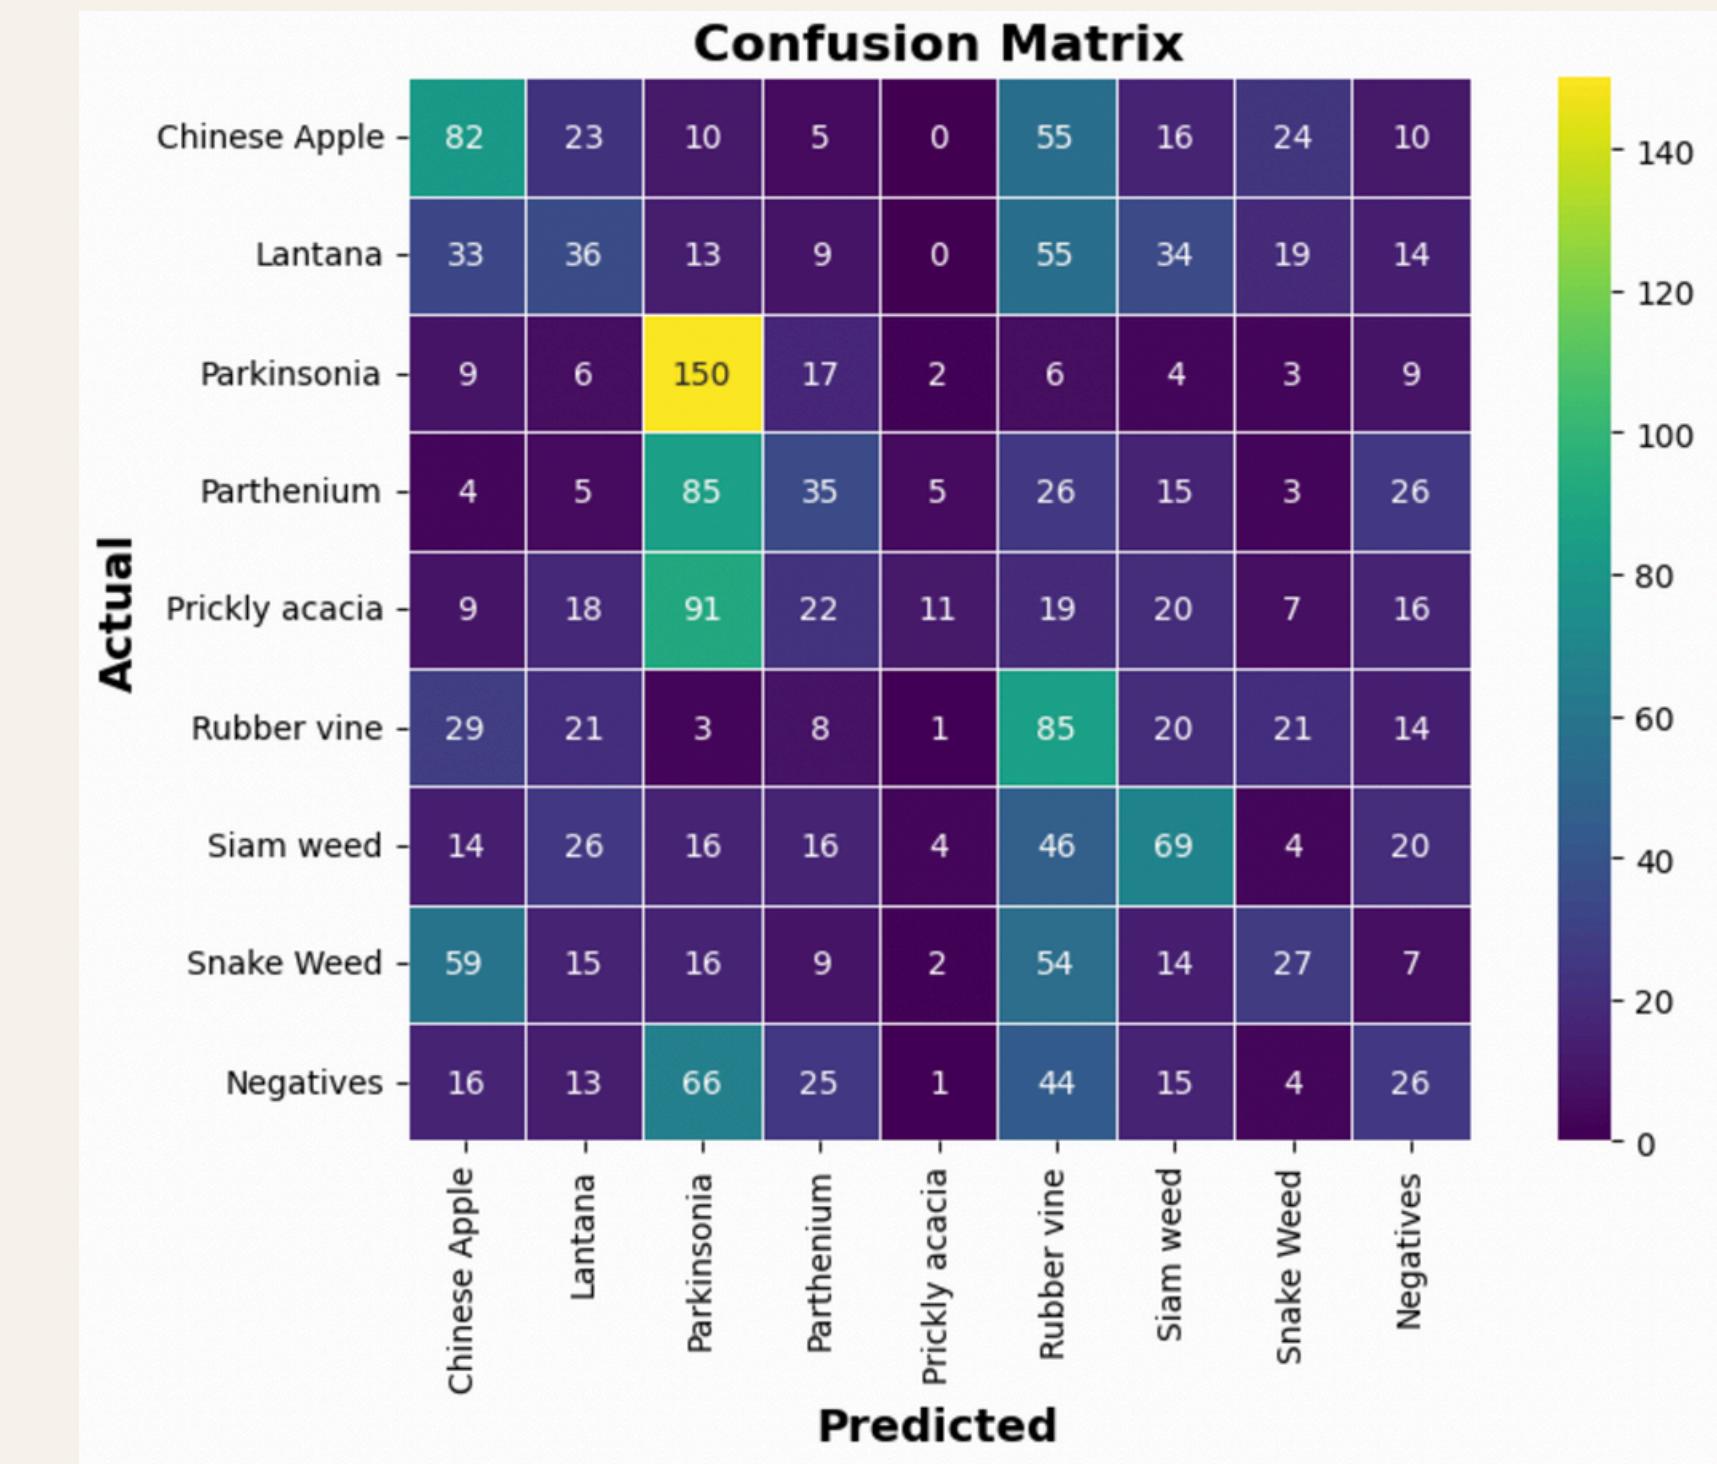
Micro-averaged Precision: 0.20

Micro-averaged Recall: 0.20

Micro-averaged F1 score: 0.20

Classification Report for Random Forest:

	precision	recall	f1-score	support
Chinese Apple	0.26	0.29	0.27	225
Lantana	0.15	0.15	0.15	213
Parkinsonia	0.16	0.72	0.26	206
Parthenium	0.11	0.18	0.14	204
Prickly acacia	0.18	0.03	0.05	213
Rubber vine	0.11	0.48	0.18	202
Siam weed	0.22	0.24	0.23	215
Snake Weed	0.15	0.22	0.17	203
Negatives	0.66	0.12	0.21	1821
accuracy			0.20	3502
macro avg	0.22	0.27	0.18	3502
weighted avg	0.43	0.20	0.20	3502



# Support Vector Machine (SVM)

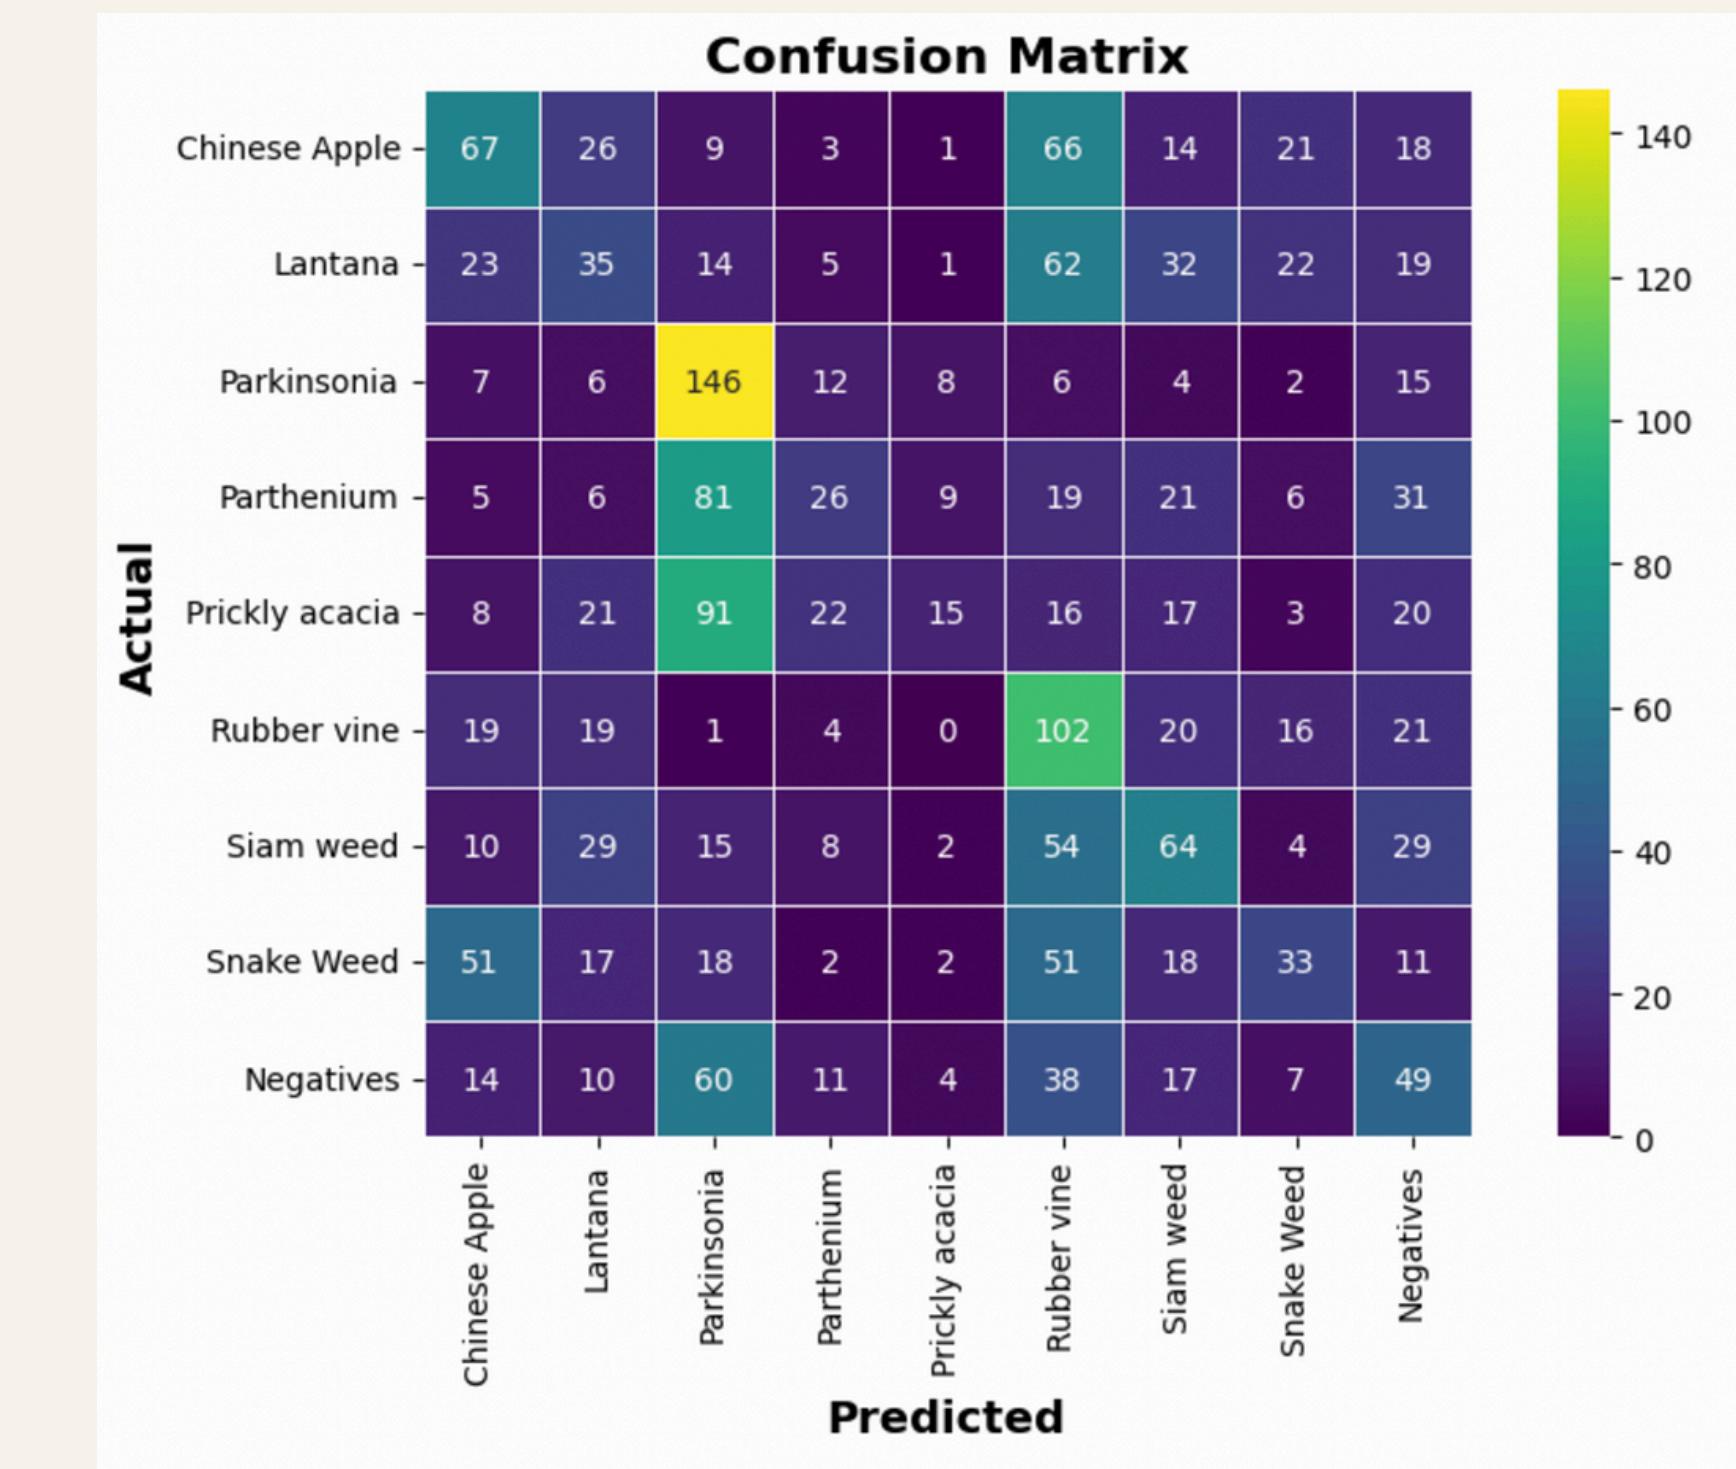
Micro-averaged Precision: 0.26

Micro-averaged Recall: 0.26

Micro-averaged F1 score: 0.26

Classification Report for SVM:

	precision	recall	f1-score	support
Chinese Apple	0.25	0.26	0.25	225
Lantana	0.18	0.19	0.19	213
Parkinsonia	0.17	0.68	0.27	206
Parthenium	0.17	0.11	0.14	204
Prickly acacia	0.19	0.03	0.05	213
Rubber vine	0.14	0.55	0.23	202
Siam weed	0.18	0.33	0.23	215
Snake Weed	0.13	0.12	0.12	203
Negatives	0.68	0.25	0.36	1821
accuracy			0.26	3502
macro avg	0.23	0.28	0.20	3502
weighted avg	0.44	0.26	0.28	3502



# *Evaluation of linear models*

## *Micro-Averaged*

	Precision	Recall	F1
kNN	0.22	0.22	0.22
Random Forest	0.20	0.20	0.20
SVM	0.26	0.26	0.26

# 04

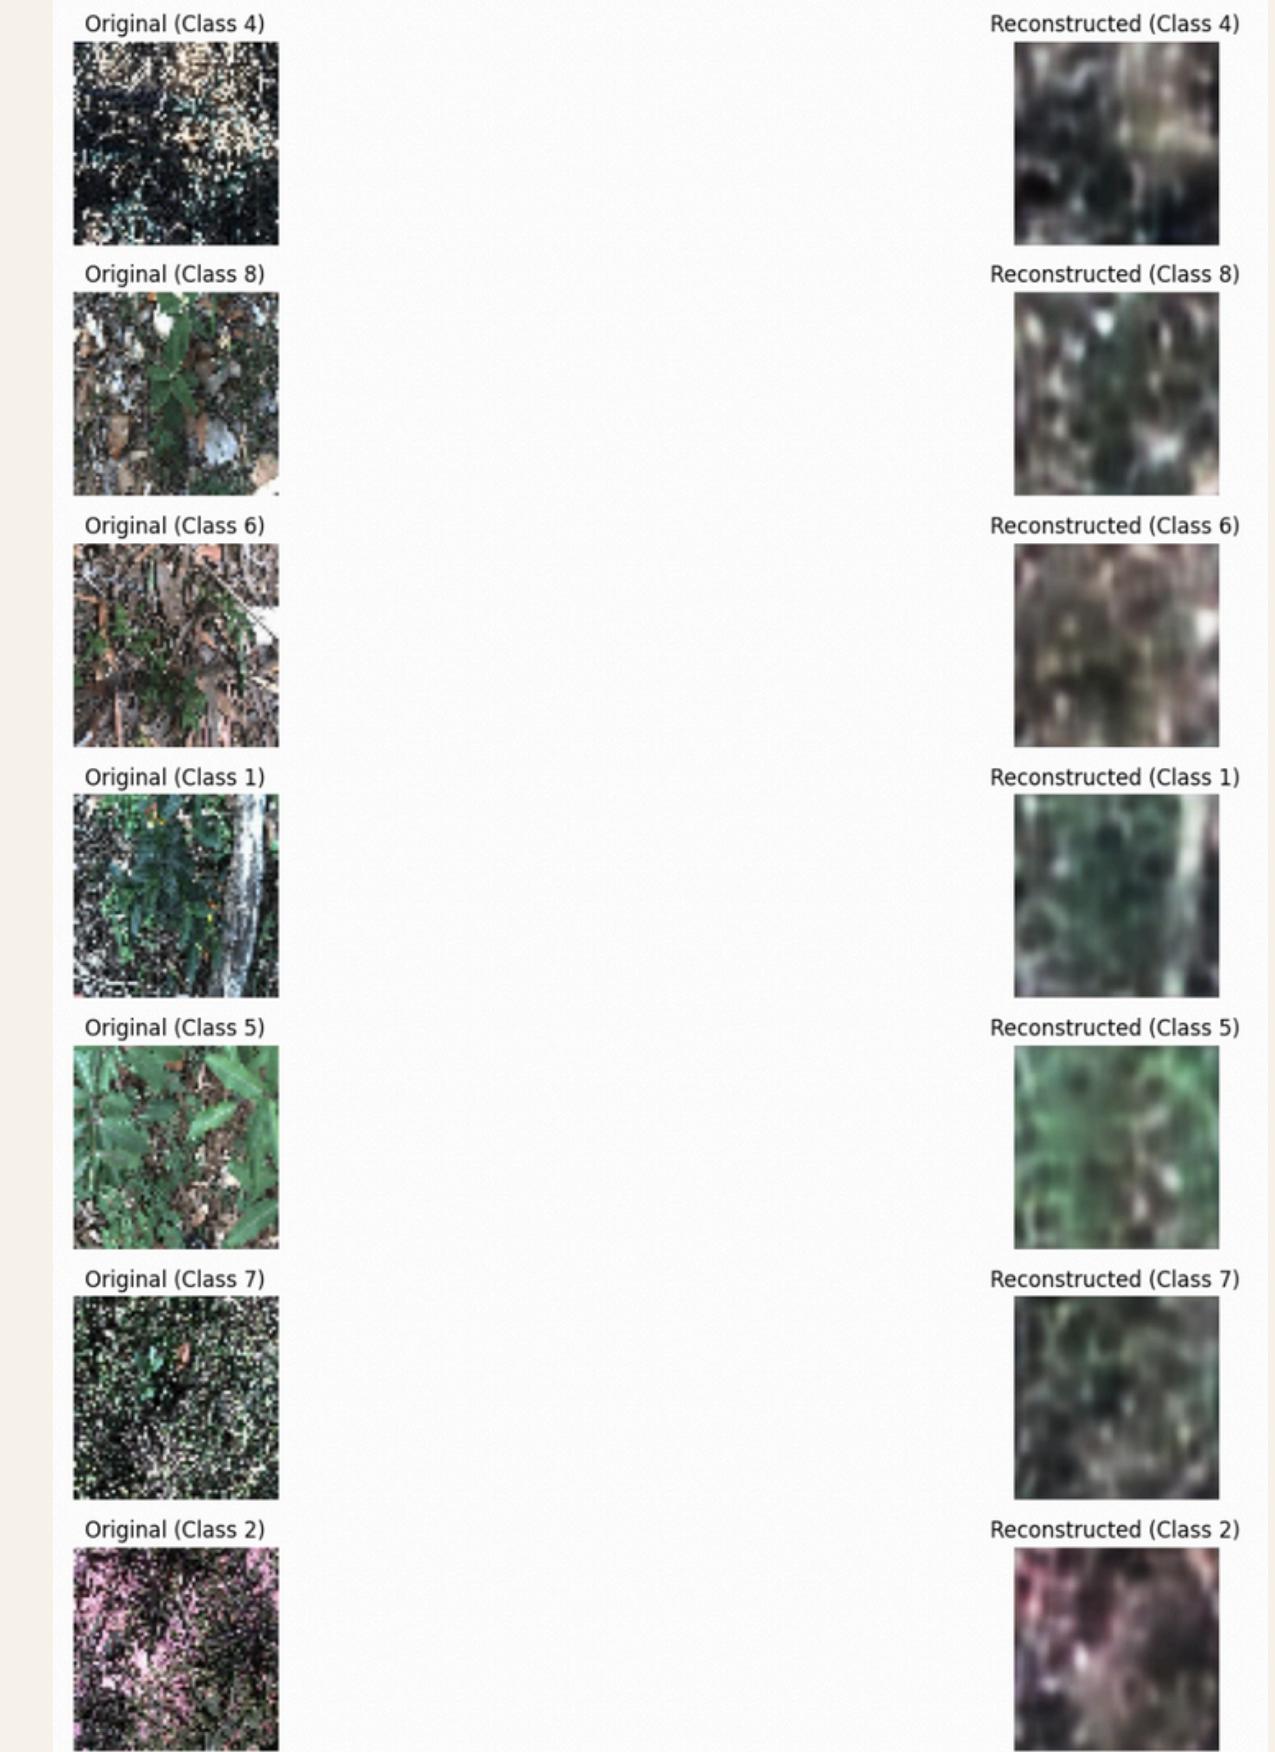
## Linear Models with Autoencoders

# *Autoencoders with linear models*

Employed autoencoder to extract more meaningful features

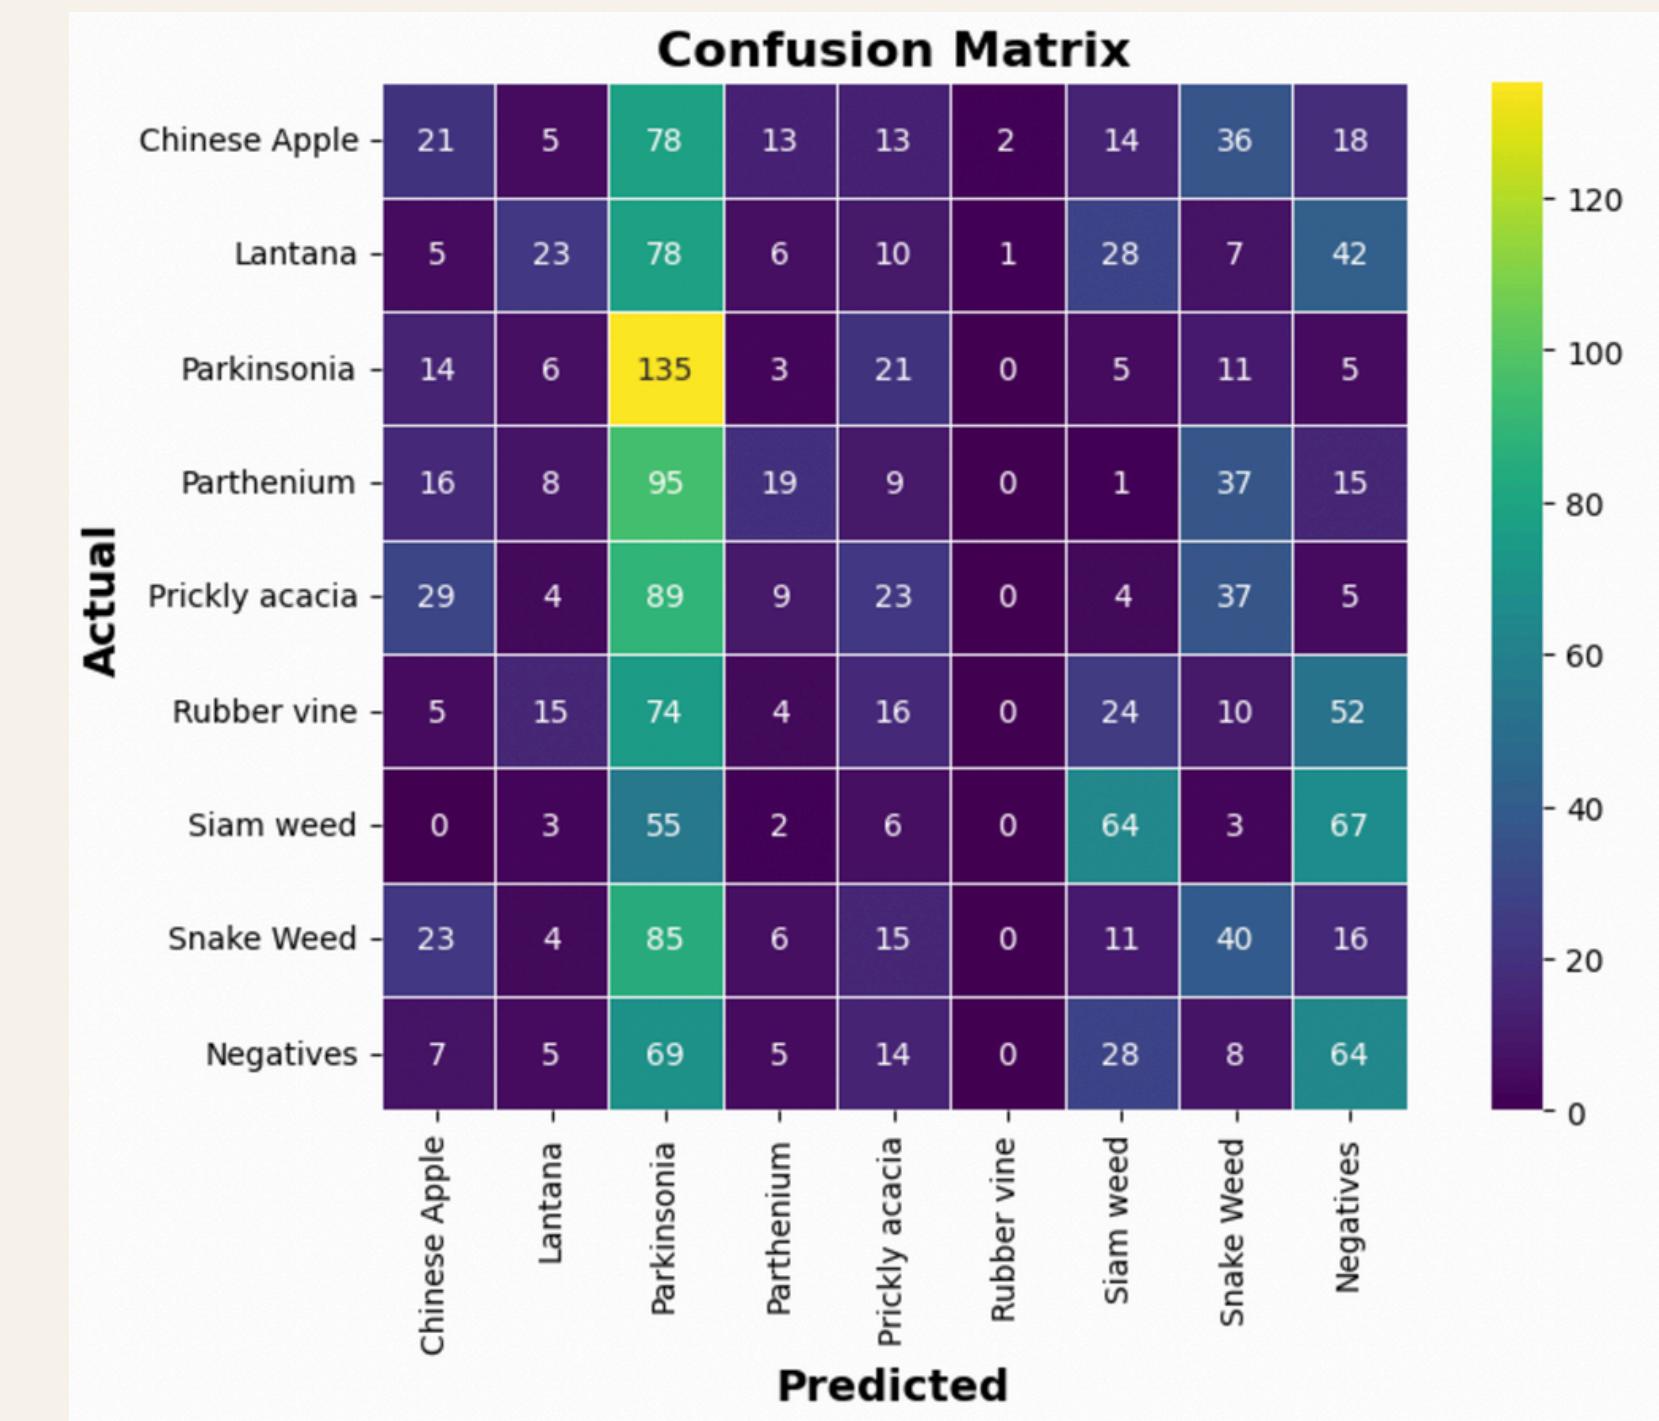
Operated in two stages:

1. compresses images into lower-dimensional
2. Decoder constructs images from the compressed version



# *kNN using autoencoder*

Macro-averaged Precision:	0.21			
Macro-averaged Recall:	0.22			
Macro-averaged F1 score:	0.19			
Classification Report for kNN:				
	precision	recall	f1-score	support
Chinese Apple	0.17	0.10	0.13	200
Lantana	0.32	0.12	0.17	200
Parkinsonia	0.18	0.68	0.28	200
Parthenium	0.28	0.10	0.14	200
Prickly acacia	0.18	0.12	0.14	200
Rubber vine	0.00	0.00	0.00	200
Siam weed	0.36	0.32	0.34	200
Snake Weed	0.21	0.20	0.21	200
Negatives	0.23	0.32	0.26	200
accuracy			0.22	1800
macro avg	0.21	0.22	0.19	1800
weighted avg	0.21	0.22	0.19	1800



# Random Forest using autoencoder

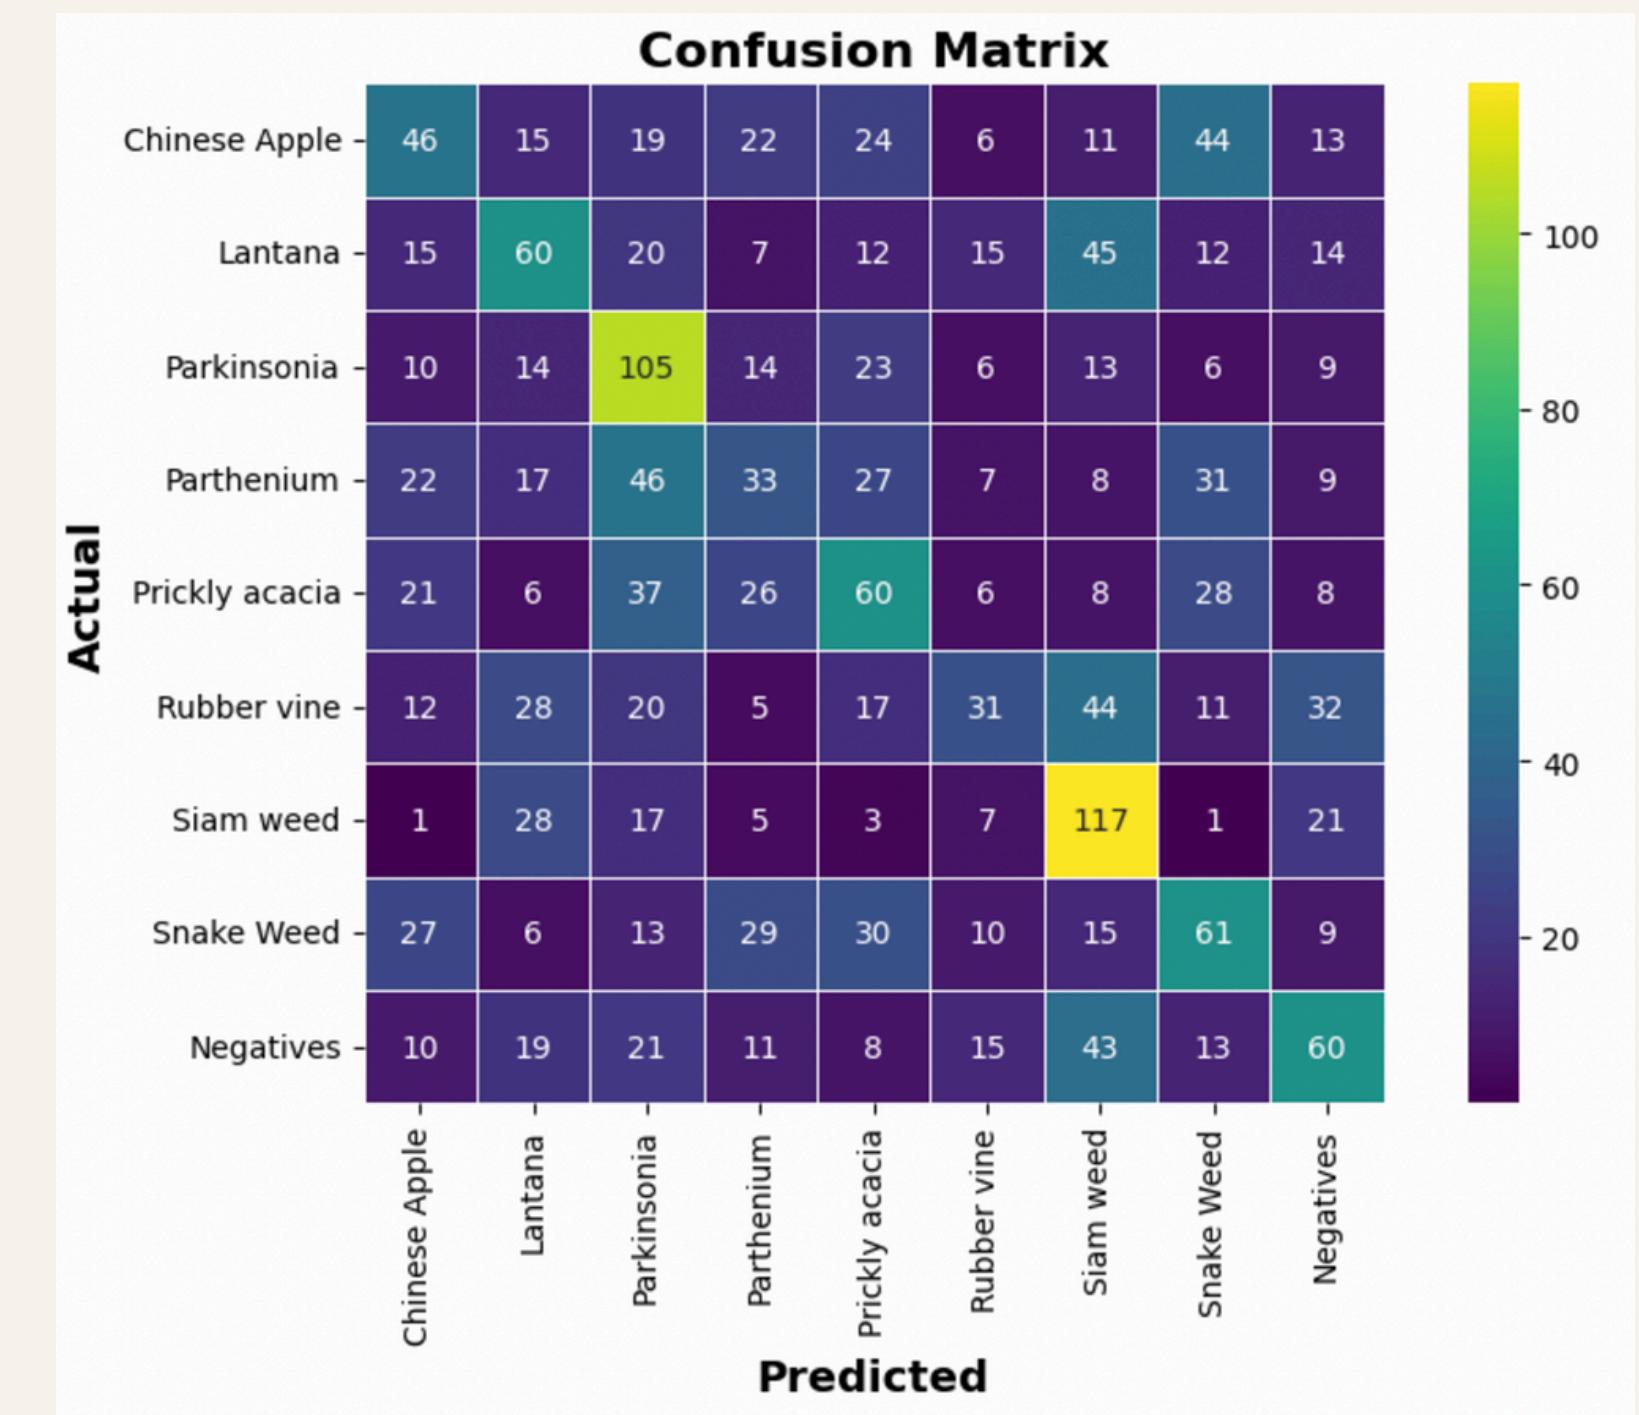
Macro-averaged Precision: 0.31

Macro-averaged Recall: 0.32

Macro-averaged F1 score: 0.31

Classification Report for Random Forest:

	precision	recall	f1-score	support
Chinese Apple	0.28	0.23	0.25	200
Lantana	0.31	0.30	0.31	200
Parkinsonia	0.35	0.53	0.42	200
Parthenium	0.22	0.17	0.19	200
Prickly acacia	0.29	0.30	0.30	200
Rubber vine	0.30	0.15	0.20	200
Siam weed	0.38	0.58	0.46	200
Snake Weed	0.29	0.30	0.30	200
Negatives	0.34	0.30	0.32	200
accuracy			0.32	1800
macro avg	0.31	0.32	0.31	1800
weighted avg	0.31	0.32	0.31	1800

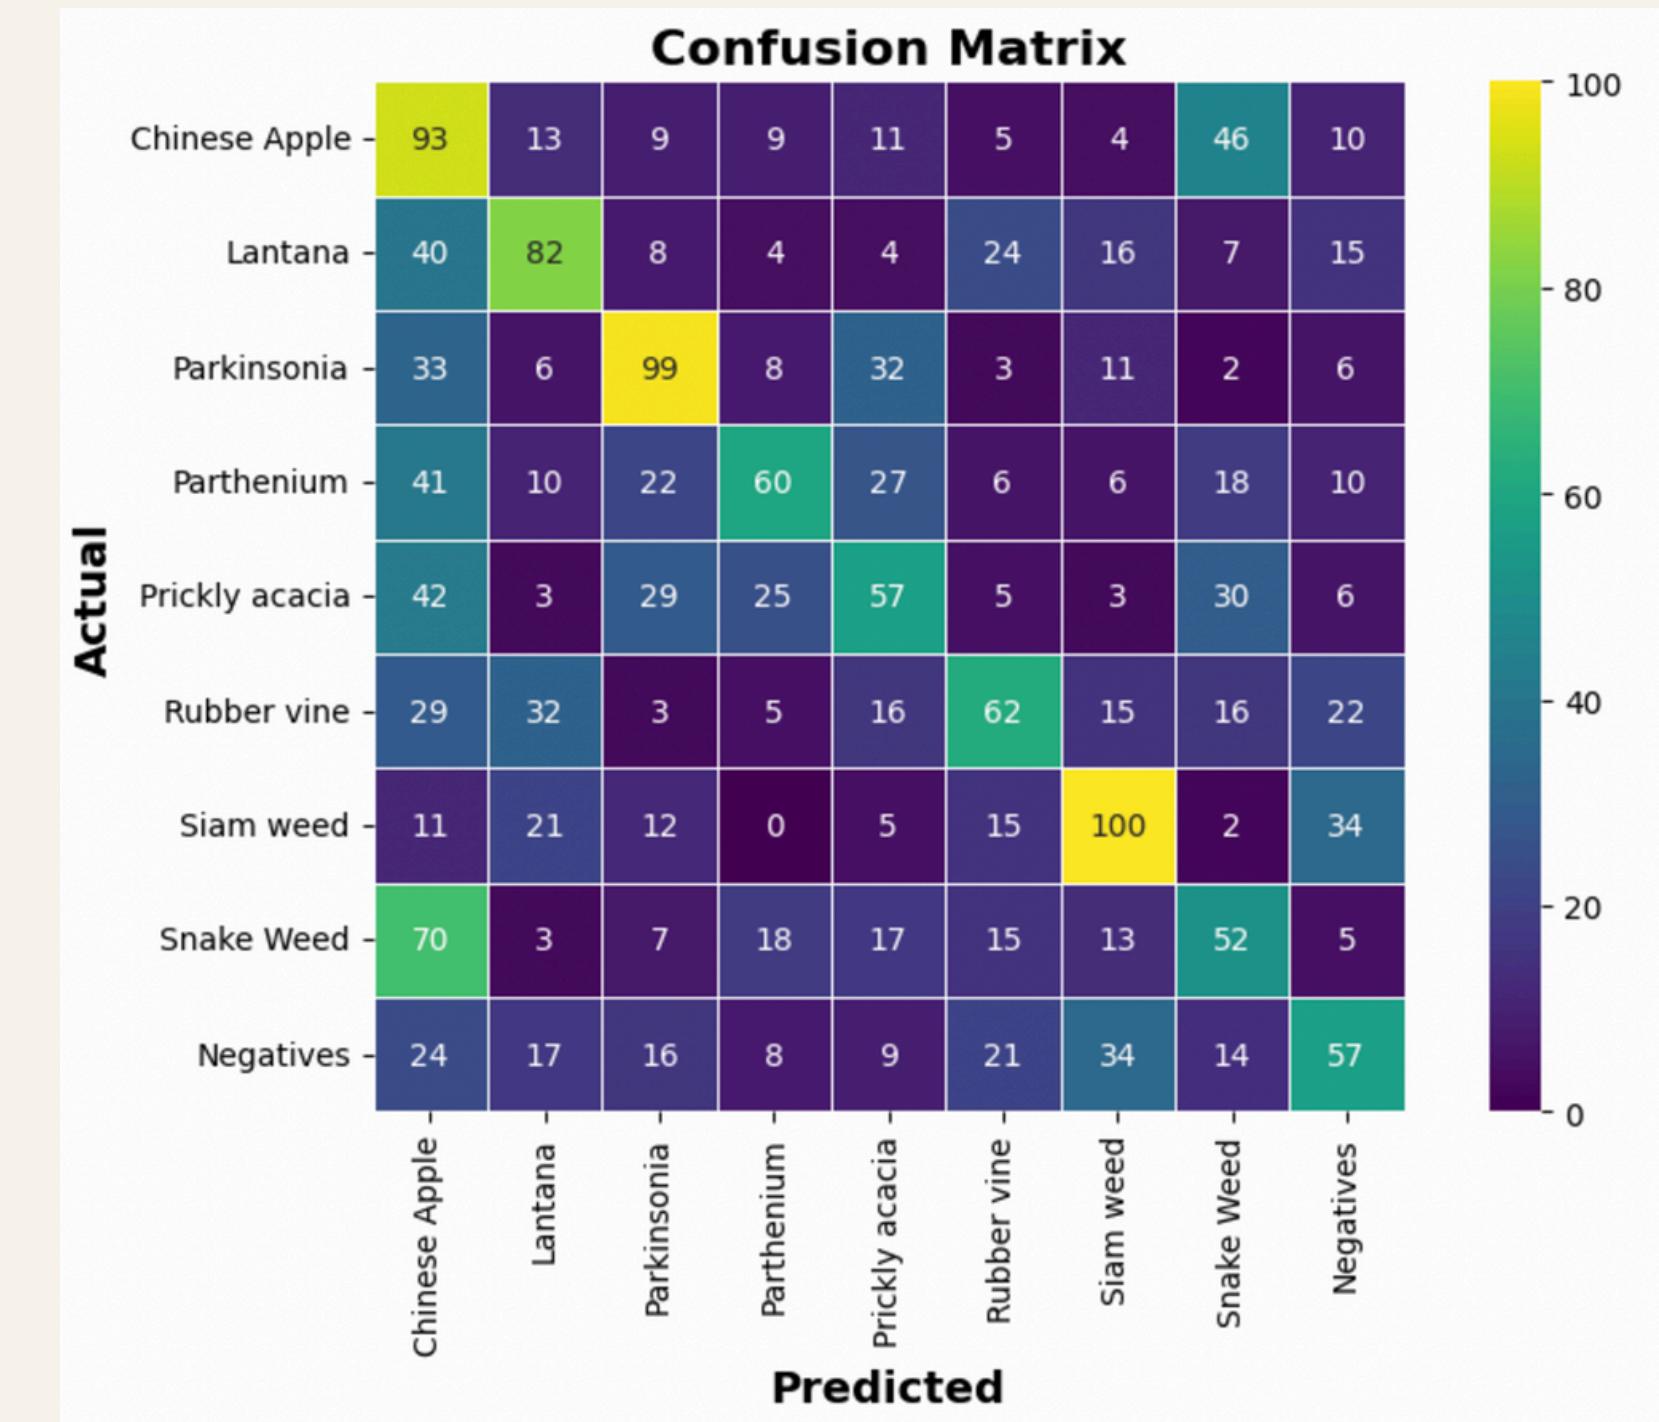


# *SVM using autoencoder*

```

Macro-averaged Precision: 0.38
Macro-averaged Recall: 0.37
Macro-averaged F1 score: 0.37
Classification Report for SVM:
precision    recall   f1-score   support
Chinese Apple      0.24      0.47      0.32      200
Lantana            0.44      0.41      0.42      200
Parkinsonia        0.48      0.49      0.49      200
Parthenium          0.44      0.30      0.36      200
Prickly acacia     0.32      0.28      0.30      200
Rubber vine         0.40      0.31      0.35      200
Siam weed           0.50      0.50      0.50      200
Snake Weed          0.28      0.26      0.27      200
Negatives           0.35      0.28      0.31      200
accuracy             -         -         -         1800
macro avg            0.38      0.37      0.37      1800
weighted avg         0.38      0.37      0.37      1800

```



# 05

## CNN and GradCAM

# Convolutional Neural Network (CNN)

```
# Define the model
inputs = Input(shape=(224, 224, 4))

x = layers.Conv2D(64, (3, 3), padding="same", strides=2)(inputs)
x = layers.Conv2D(64, (3, 3), padding="same", strides=2)(x)
x = layers.BatchNormalization()(x)
layer1 = layers.ReLU()(x)

x = layers.Conv2D(64, (3, 3), padding="same")(layer1)
x = layers.Conv2D(64, (3, 3), padding="same")(x)
x = layers.BatchNormalization()(x)
layer2 = layers.ReLU()(x)

x = layers.Conv2D(128, (3, 3), padding="same")(layer2)
x = layers.Conv2D(128, (3, 3), padding="same")(x)
x = layers.BatchNormalization()(x)
layer3 = layers.ReLU()(x)

x = layers.Conv2D(128, (3, 3), padding="same")(layer3)
x = layers.Conv2D(128, (3, 3), padding="same")(x)
x = layers.BatchNormalization()(x)
layer4 = layers.ReLU()(x)

x = layers.Conv2D(256, (3, 3), padding="same", strides=2)(layer4)
x = layers.Conv2D(256, (3, 3), padding="same", strides=2)(x)
x = layers.BatchNormalization()(x)
layer5 = layers.ReLU()(x)

x = layers.GlobalMaxPooling2D()(x)
x = layers.Dropout(0.5)(x)
outputs = layers.Dense(9, activation="softmax")(x)
```

```
Epoch 1/20  
197/197 [=====] - 738s 4s/step - loss: 2.5979 - accuracy: 0.2319 - precision: 0.2587 - recall: 0.1178 - val_loss: 4.6115 - val_accuracy: 0.1211 - val_precision: 0.1278 - val_recall: 0.1200  
Epoch 2/20  
197/197 [=====] - 685s 3s/step - loss: 1.8098 - accuracy: 0.3300 - precision: 0.4214 - recall: 0.1208 - val_loss: 2.6872 - val_accuracy: 0.2422 - val_precision: 0.3079 - val_recall: 0.1656  
Epoch 3/20  
197/197 [=====] - 671s 3s/step - loss: 1.6615 - accuracy: 0.3732 - precision: 0.5052 - recall: 0.1532 - val_loss: 1.7049 - val_accuracy: 0.3711 - val_precision: 0.4748 - val_recall: 0.1467  
Epoch 4/20  
197/197 [=====] - 682s 3s/step - loss: 1.5899 - accuracy: 0.4068 - precision: 0.5437 - recall: 0.1897 - val_loss: 1.8303 - val_accuracy: 0.3711 - val_precision: 0.4619 - val_recall: 0.2356  
Epoch 5/20  
197/197 [=====] - 645s 3s/step - loss: 1.5404 - accuracy: 0.4243 - precision: 0.5645 - recall: 0.1986 - val_loss: 1.6364 - val_accuracy: 0.4011 - val_precision: 0.5466 - val_recall: 0.2344
```



```
Epoch 15/20  
197/197 [=====] - 636s 3s/step - loss: 1.1302 - accuracy: 0.5924 - precision: 0.7173 - recall: 0.4433 - val_loss: 1.5581 - val_accuracy: 0.4478 - val_precision: 0.6359 - val_recall: 0.2833  
Epoch 16/20  
197/197 [=====] - 653s 3s/step - loss: 1.0938 - accuracy: 0.6052 - precision: 0.7321 - recall: 0.4637 - val_loss: 1.4341 - val_accuracy: 0.5589 - val_precision: 0.6026 - val_recall: 0.4633  
Epoch 17/20  
197/197 [=====] - 24642s 126s/step - loss: 1.0498 - accuracy: 0.6192 - precision: 0.7401 - recall: 0.4895 - val_loss: 3.1679 - val_accuracy: 0.2600 - val_precision: 0.2845 - val_recall: 0.2200  
Epoch 18/20  
197/197 [=====] - 649s 3s/step - loss: 1.0525 - accuracy: 0.6219 - precision: 0.7330 - recall: 0.4859 - val_loss: 1.2306 - val_accuracy: 0.5678 - val_precision: 0.6558 - val_recall: 0.5144  
Epoch 19/20  
197/197 [=====] - 611s 3s/step - loss: 1.0122 - accuracy: 0.6378 - precision: 0.7526 - recall: 0.5065 - val_loss: 1.0664 - val_accuracy: 0.6167 - val_precision: 0.7210 - val_recall: 0.4767  
Epoch 20/20  
197/197 [=====] - 640s 3s/step - loss: 0.9935 - accuracy: 0.6398 - precision: 0.7514 - recall: 0.5149 - val_loss: 1.1733 - val_accuracy: 0.5922 - val_precision: 0.7003 - val_recall: 0.5089
```

## 1st Epoch

### **validation set**

- **precision: 0.1278**
- **recall: 0.1200**

## Last 20th Epoch

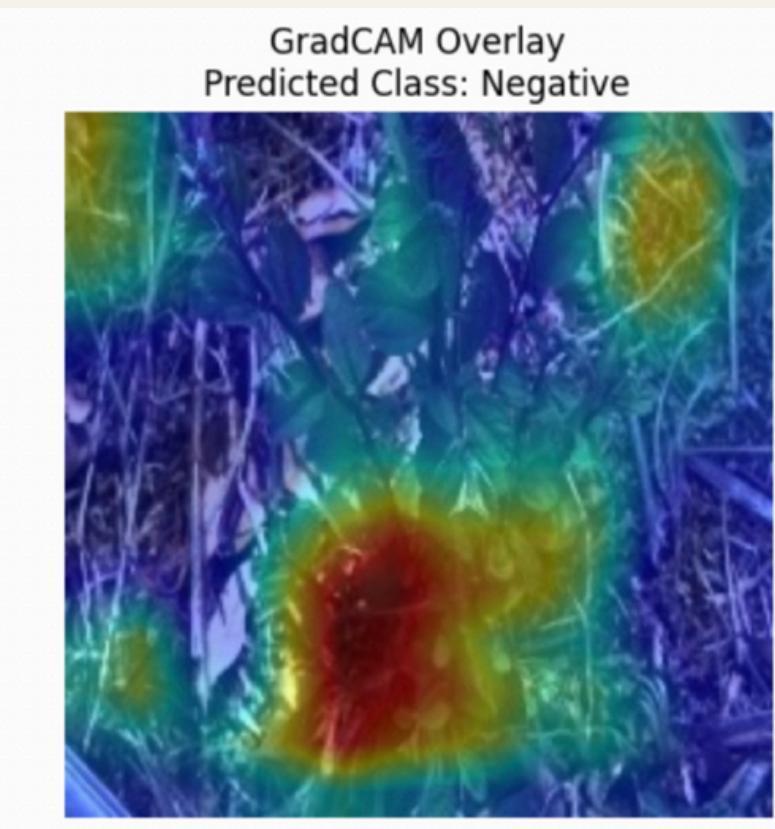
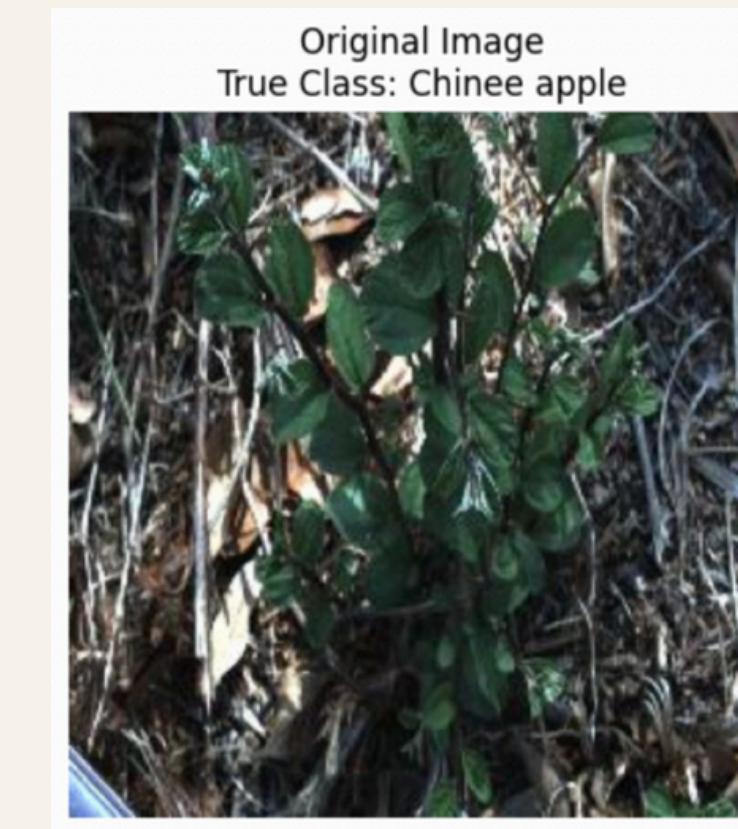
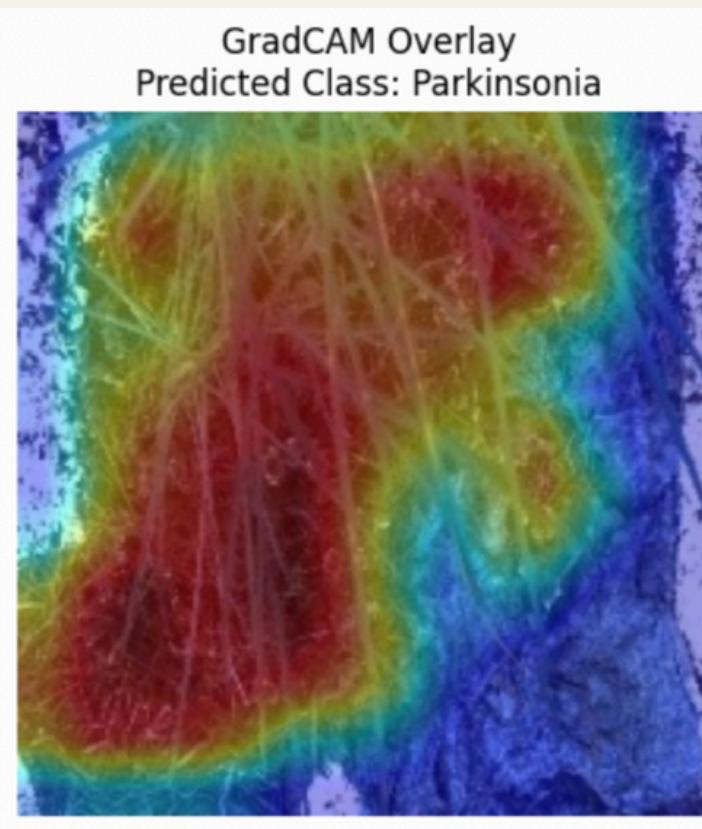
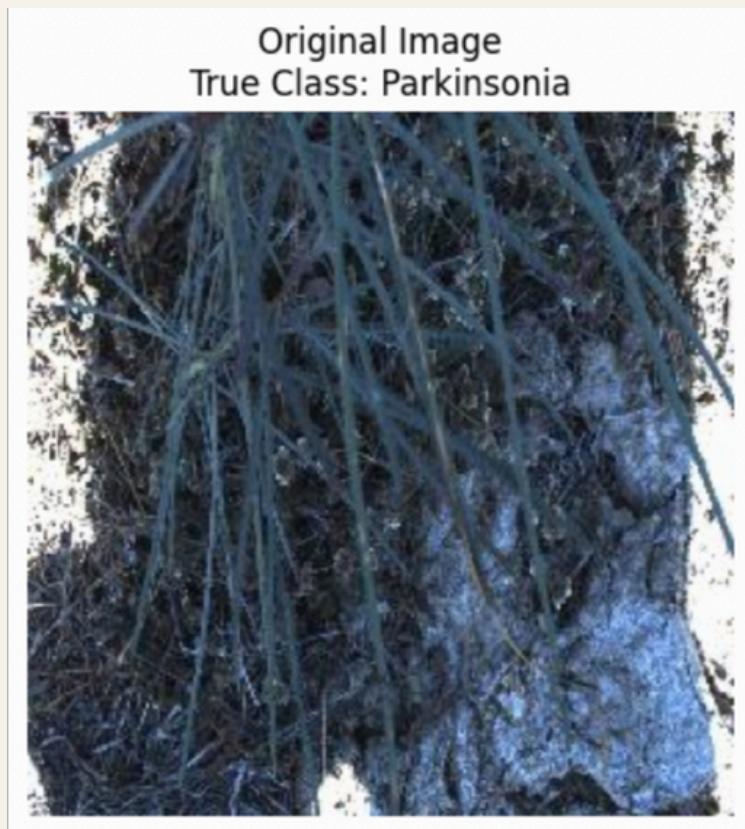
### **validation set**

- **precision: 0.7003**
- **recall: 0.5089**

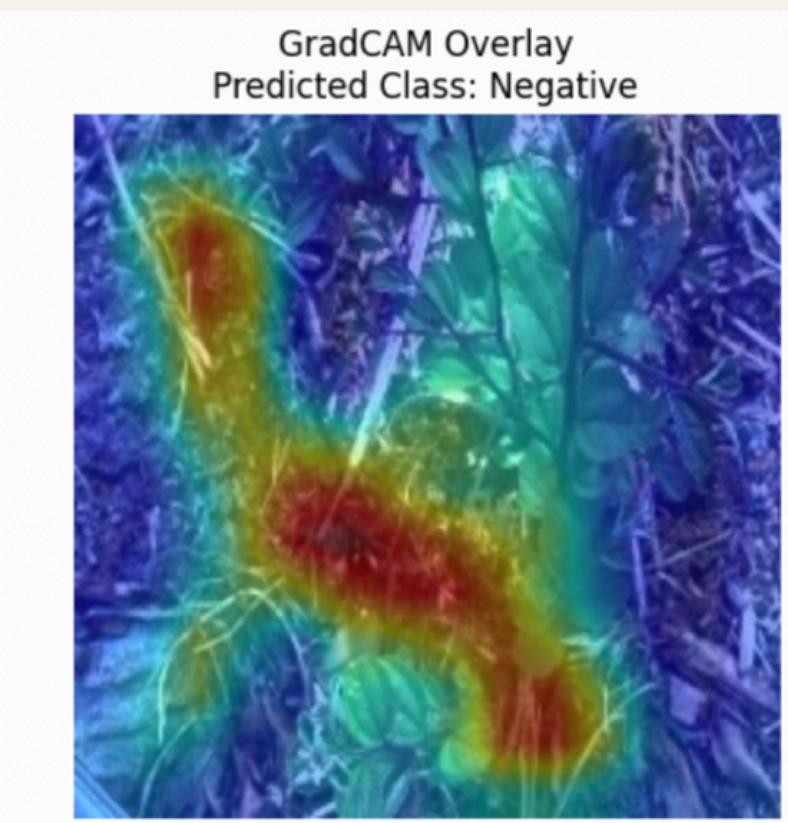
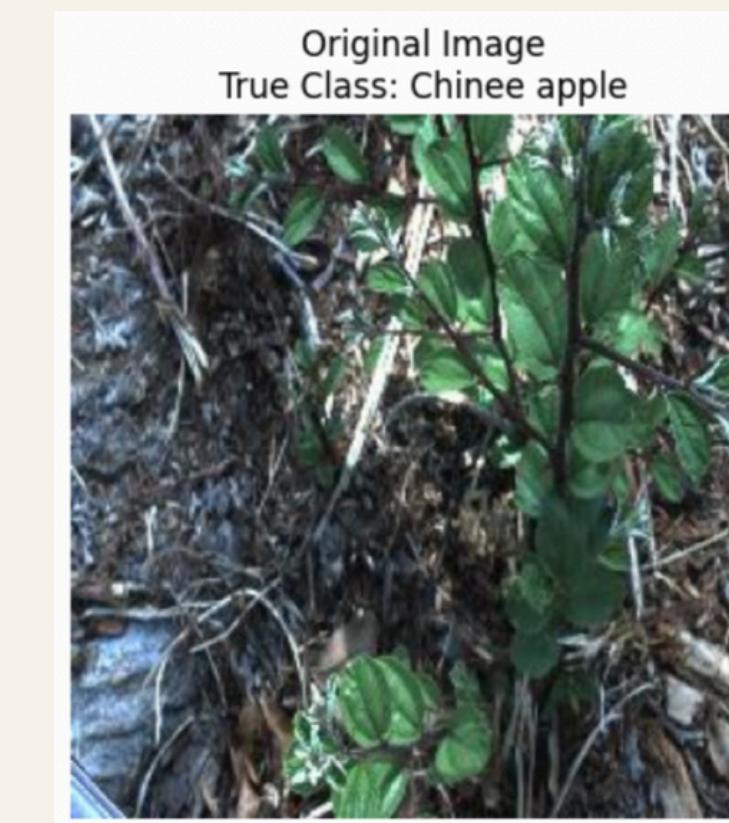
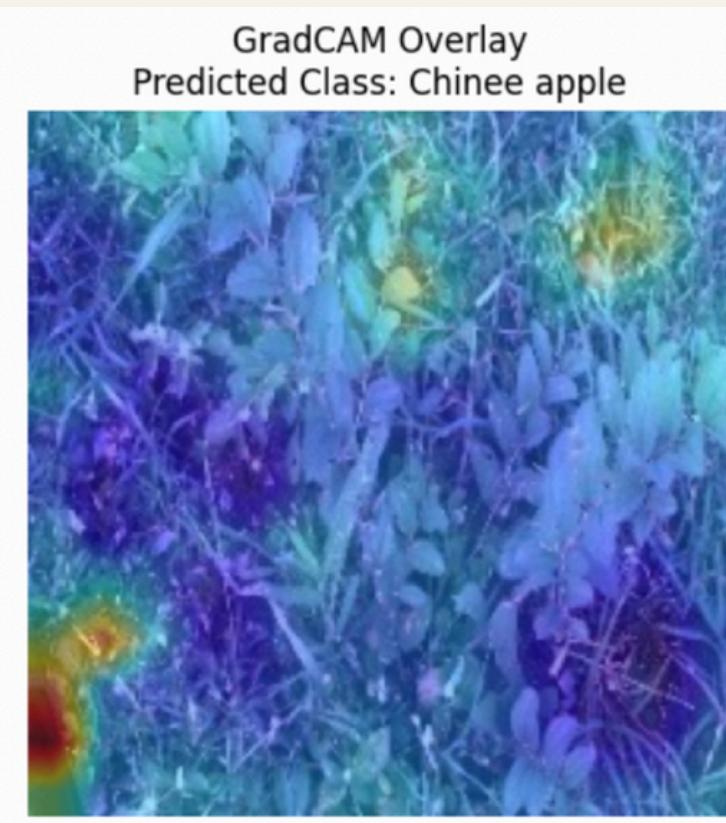
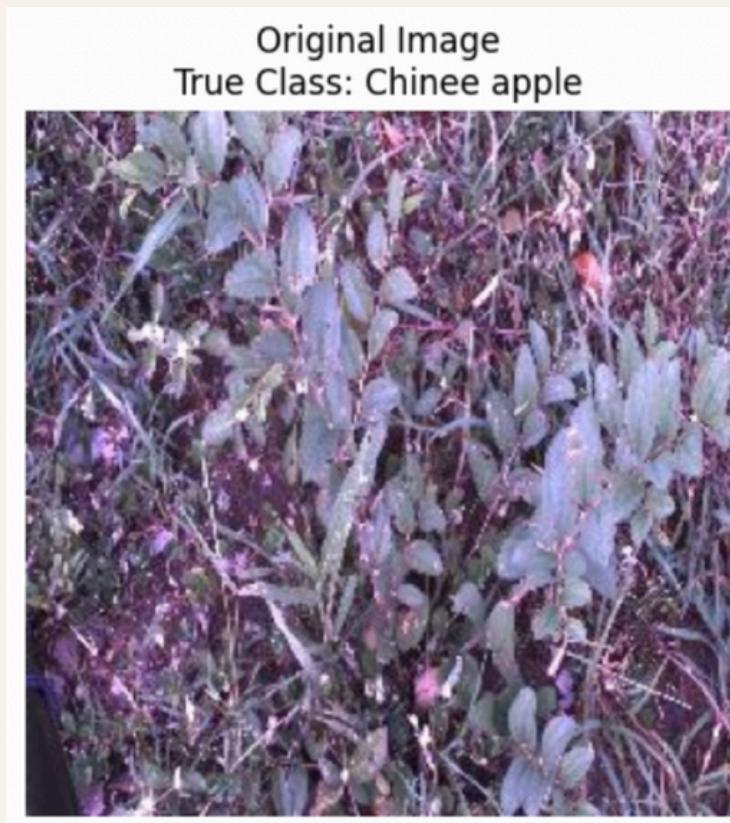
# *Convolutional Neural Network (CNN)*

Micro-averaged Precision:	0.62			
Micro-averaged Recall:	0.62			
Micro-averaged F1 score:	0.62			
precision recall f1-score support				
Chinese Apple	0.75	0.04	0.08	225
Lantana	0.37	0.77	0.50	213
Parkinsonia	0.56	0.67	0.61	206
Parthenium	0.41	0.75	0.53	204
Prickly acacia	0.63	0.75	0.69	213
Rubber vine	0.45	0.83	0.59	202
Siam weed	0.50	0.77	0.60	215
Snake Weed	0.41	0.52	0.46	203
Negatives	0.90	0.60	0.72	1821
accuracy			0.62	3502
macro avg	0.55	0.63	0.53	3502
weighted avg	0.72	0.62	0.62	3502

# *GradCAM (correctly identified)*



# *GradCAM (wrongly identified)*



*06*

## Conclusion

# *Conclusion*

*01.*

This project focused on improving weed classification in Australian agriculture using machine learning, overcoming challenges like class imbalance.

*02.*

While traditional models had limitations, Convolutional Neural Networks (CNNs) outperformed them by capturing complex patterns.

*03.*

Grad-CAM insights confirmed the CNN's effectiveness, offering a transformative solution for weed management in farming.

*Thank you*